

BERI

Bluespec Extensible RISC Implementation

Hardware Reference

Version 1.2

This interim document is not released for public consumption

Robert N. M. Watson, David Chisnall, Brooks Davis, Wojciech Koszek,
Simon W. Moore, Steven J. Murdoch, Peter G. Neumann, Robert Norton,
Michael Roe, Jonathan Woodruff
SRI International and the University of Cambridge¹

March 5, 2014

¹Sponsored by the Defense Advanced Research Projects Agency (DARPA) and the Air Force Research Laboratory (AFRL), under contract FA8750-10-C-0237. The views, opinions, and/or findings contained in this report are those of the authors and should not be interpreted as representing the official views or policies, either expressed or implied, of the Defense Advanced Research Projects Agency or the Department of Defense.

Contents

1	Introduction	7
1.1	Bluespec Extensible RISC Implementation (BERI)	7
1.2	BERI and BERI2 Prototypes	7
1.3	Capability Hardware Enhanced RISC Instructions (CHERI)	8
1.4	Getting BERI	8
1.5	Using BERI	8
1.6	Licensing	8
1.7	Version History	9
1.8	Document Structure	9
2	The BERI ISA	11
2.1	BERI CPU Features and ISA	11
2.1.1	MIPS Instructions	12
2.1.2	Limitations to Coprocessor 0 Support	21
2.1.3	Modifications to the MIPS TLB Model	21
2.1.4	Limitations to Generic Coprocessor Support	22
2.1.5	The BERI Floating-Point Unit (FPU)	22
2.1.6	Selected Additions from Later MIPS ISA Versions	23
2.1.7	Virtual Address Space	23
2.2	CHERI ISA Extensions	26
2.3	Other BERI ISA Extensions	26
2.4	BERI2 Specific Features	28
2.4.1	Multi-Threading	28
3	The BERI Programmable Interrupt Controller	29
3.1	Sources	29
3.2	Source Numbers and Base Addresses on BERI	29
3.3	Config Registers: PIC_CONFIG_X	30
3.4	Interrupt Pending Bits	30
3.5	Reset State	32
3.6	Safe Handling of Interrupts	32

4	The BERI Processor Implementation	33
4.1	Bluespec	33
4.2	Directory Layout	33
4.3	Key Files	34
4.4	Conditionally Compiled Features	34
5	Simulating BERI	37
5.1	Software Dependencies to Build BERI	37
5.2	Checking out BERI	39
5.3	Building the BERI Simulator	39
5.4	Configuring the BERI Simulator	40
5.5	Simulating BERI	40
5.6	Running the BERI test suite	42
6	Using the BERI Debug Unit	43
6.1	Communicating with the BERI Debug Unit	43
6.2	BERI Debug Registers	43
6.3	BERI Debug Instructions	44
6.4	BERI Debug Responses	45
7	The BERI Unit-Test Suite	47
7.1	The BERI Unit-Test Environment	47
7.2	Software Dependencies	47
7.2.1	BERI Test-Suite Directory Layout	48
7.2.2	BERI ISA Extensions for Testing	48
7.2.3	Unit Test Support Library	48
7.3	Running the BERI Test Suite	49
7.3.1	Jenkins	50
7.4	Unit-Test Structure	50
7.4.1	Test Types	50
7.4.2	Test Structure	50
7.4.3	Test Termination	51
7.4.4	Connecting New Tests to the Build	51
7.4.5	Test Attributes	51
7.5	Example Unit Test: Register Zero	51
7.5.1	Register Zero Test Code	53
7.5.2	Register Zero Nose Assertions	53
7.6	Conclusion	53
8	BERI on Altera FPGAs	57
8.1	Building BERI for Synthesis	57
8.2	The Altera Development Environment	57
8.3	Synthesizing BERI	60

9	BERI on Terasic tPad and DE4	61
9.1	BERI Configuration on the tPad and DE4	61
9.1.1	Physical Address Space on the tPad	61
9.1.2	Physical Address Space on the DE4	61
9.2	Altera IP Cores	61
9.3	Cambridge IP Cores	64
9.3.1	The DE4 Multitouch LCD	64
9.3.2	HDMI Chip Configuration via I2C	67
9.4	Standalone HDMI Output	68
9.4.1	Reconfigurable Video Pixel Clock	68
9.4.2	HDMI Timing Driver	69
9.5	Temperature and fan control	69
9.6	Bluespec 16550 UART	69
9.7	Terasic hard peripherals	70
9.7.1	Intel StrataFlash 64M NOR flash	70

Abstract

This document is the *BERI Hardware Reference* documenting the BERI processor developed by SRI International and the University of Cambridge. The *Hardware reference* is targeted at hardware and software developers working with the BERI prototype in simulation and synthesized to FPGA targets. The manual includes chapters on the implementation status of the 64-bit MIPS and CHERI ISAs in the prototype, the BERI Programmable Interrupt Controller (PIC), the BERI processor implementation, the BERI processor simulation, the BERI debug unit, the BERI test suite, and information on using BERI with Altera FPGAs and Terasic tPad and DE4 boards. The companion *CHERI User's Guide* and *BERI Software Reference* provide operational and software development information for the platform.

Acknowledgments

The authors of this report thank other members of the CTSRD team, and our past and current research collaborators at SRI and Cambridge:

Ross J. Anderson	Jonathan Anderson	Gregory Chadwick	Nirav Dave
Khilan Gudka	Jong Hun Han	Alex Horsman	Alexandre Joannou
Asif Khan	Myron King	Ben Laurie	Patrick Lincoln
Anil Madhavapeddy	Ilias Marinos	Ed Maste	Andrew Moore
Will Morland	Alan Mujumdar	Prashanth Mundkur	Philip Paeps
Colin Rothwell	John Rushby	Hassen Saidi	Muhammad Shahbaz
Stacey Son	Richard Uhler	Philip Withnall	Bjoern Zeeb

The CTSRD team wishes to thank its external oversight group for significant support and contributions:

Lee Badger	Simon Cooper	Rance DeLong	Jeremy Epstein
Virgil Gligor	Li Gong	Mike Gordon	Steven Hand
Andrew Herbert	Warren A. Hunt Jr.	Doug Maughan	Greg Morrisett
Brian Randell	Kenneth F. Shottling	Joe Stoy	Tom Van Vleck
Samuel M. Weber			

Finally, we are grateful to Howie Shrobe, MIT professor and past DARPA CRASH program manager, who has offered both technical insight and support throughout this work.

Chapter 1

Introduction

This is the *BERI Hardware Reference* for the Bluespec Extensible RISC Implementation (BERI) prototype. This document describes the status of the processor prototype, early information on the implementation itself, and is a reference for various aspects of the hardware platform, such as the BERI Programmable Interrupt Controller (PIC) and supported Altera, Terasic, and Cambridge/SRI IP cores. It complements the *BERI Software Reference*, which describes the BERI software development environment, and the *CHERI User's Guide* which discusses CHERI extensions to the BERI software environment.

1.1 Bluespec Extensible RISC Implementation (BERI)

The Bluespec Extensible RISC Implementation (BERI) is a platform for performing research into the hardware-software interface that has been developed as part of the CTSRD project. It consists of a 64-bit MIPS CPU implemented in Bluespec System Verilog and a complete software stack based on FreeBSD, Clang/LLVM, and a range of popular open-source software products such as Apache and Chromium. Wherever possible, BERI makes use of BSD- and Apache-licensed software to maximize opportunities for technology transition.

1.2 BERI and BERI2 Prototypes

The first BERI prototype was developed between 2010-2012 by Jonathan Woodruff, based on the 64-bit MIPS ISA implementation originally created by Gregory Chadwick. We are also developing multi-core support for BERI1. That version is the primary focus of the current *BERI Hardware Reference* and *CHERI User's Guide*.

BERI2 is the second version of the prototype developed between 2011-2012 by Nirave Dave and Robert Norton using a stylized form of Bluespec in order to better support formal verification. BERI2 is designed to support multi-threaded as well as multi-core operation.

Although BERI and BERI2 share significant infrastructure (for example, memory subsystems and simulated peripheral buses), we do not currently envision convergence of the two implementations. Instead, we expect a gradual migration of research to BERI2 as its more advanced features

mature, while keeping “production” development of BERI features on the first prototype in order to avoid premature imposition of strong stability requirements on a highly experimental implementation. In the longer term, BERI2’s support for formal methods tools should lead to much greater correctness and reliability.

1.3 Capability Hardware Enhanced RISC Instructions (CHERI)

The first major research project to be implemented on BERI is Capability Hardware Enhanced RISC Instructions (CHERI): extensions to the MIPS ISA to support fine-grained memory protection and scalable protection-domain transition within conventional MMU-based address spaces. BERI and BERI2 both include optionally compiled implementations of the CHERI ISA. An extended version of FreeBSD/BERI supporting the CHERI ISA extensions, called CheriBSD, and changes to Clang/LLVM, is available via open source, and able to use these features within UNIX applications. These are the topics of the *CHERI Architecture Document* and *CHERI User’s Guide*.

1.4 Getting BERI

We plan to distribute the CHERI prototype and reference software stack as open source via the *BERI Open-Source Downloads* page:

<http://www.cl.cam.ac.uk/research/security/ctsrd/beri-downloads.html>

At the time of writing, the prototype is not yet publicly available. It may be requested by e-mail to Dr Peter Neumann¹ at SRI International, or to Dr Robert Watson² at the University of Cambridge.

1.5 Using BERI

The BERI prototype is implemented in the Bluespec hardware description language (HDL), which may be compiled into a C-language simulator, or else synthesized for an FPGA target. The former requires access to the commercial Bluespec toolchain; the latter also requires access to the Altera FPGA toolchain.

Currently, there are two supported Altera-based FPGA boards: the Terasic DE4 and the Terasic tPad. The former offers greater FPGA performance and memory capacity; the latter includes a VGA flat panel with touch screen. We are currently in the planning phases for a port of the CHERI prototype to the NetFPGA 10G research platform, which is based on a Xilinx FPGA.

1.6 Licensing

We are currently preparing the BERI prototype for release as open source. We plan to release the hardware design, simulated peripherals, and core CPU development software tools under a version of the Apache open-source license that has been lightly modified to better take into account hardware requirements.

¹neumann@csl.sri.com

²robert.watson@cl.cam.ac.uk

We have released our extensions to the FreeBSD operating system to support BERI under a BSD license; initial support for BERI was included in FreeBSD 10.0, but further features will appear in FreeBSD 10.1. We have also released versions of FreeBSD and Clang/LLVM that support the CHERI ISA under a BSD license; these are distributed via GitHub.

We welcome contributions to the BERI project; we are only able to accept non-trivial changes when an individual or corporate contribution agreement has been signed. The BERI hardware-software license and contribution agreement may be found at:

<http://www.beri-open-systems.org/>

1.7 Version History

This is the third version of the *BERI Hardware Reference*; versions prior to 1.2 were named the *CHERI Platform Reference*.

- 1.0 The first version of the *Platform Reference Manual* was created from two relocated chapters of the then existing *User's Guide* and new content such as information on the CHERI Programmable Interrupt Controller (PIC), as well as improvements to the peripheral description (such as addition of a boot loader area to the DE4 Intel StrataFlash layout, and information on the Cambridge HDMI controller).
- 1.1 The second version is an incremental update that reflected changes in the CHERI and CHERI2 hardware platforms. Most importantly, the facilities of the new CHERI Floating Point Unit (FPU) are described. It includes documentation of higher interrupt numbers (available due to the CHERI PIC) for DE4 peripherals. Brief documentation for the Bluespec 6550 UART has been added. The new 64K L1 and L2 caches are documented. Additional divergences from the MIPS R4000 ISA are described, such as the larger 40-bit physical address space. CHERI ISA instruction information is updated. (CSBH, CSWH are no longer defined; CLLD, CSCD are now implemented; BC2F is no longer defined; CBTS, CBTU are now implemented.)
- 1.2 The *CHERI Platform Reference* has been renamed to the *BERI Hardware Reference* to reflect its focus on BERI, rather than CHERI. Test-suite attributes and simulator ISA extensions for testing are now documented; test statuses for various parts of the ISA are updated. BERI2 compilation options are now documented. Tables of floating-point instructions, conversions, and rounding modes are now included. Multicore PIC support is now documented. Further ISA extensions for core/thread identification and the thread-local storage register are documented.

1.8 Document Structure

This document is an introduction to and reference manual for the BERI processor prototype in simulation, and synthesized for Altera FPGAs on Terasic boards.

Chapter 2 enumerates portions of the MIPS and CHERI ISAs implemented by the BERI prototype. In particular, it documents sections of the MIPS ISA that have been intentionally omitted (e.g., 32-bit compatibility). It also documents the implementation status of BERI-specific ISA features, as well as aspects of the configuration of reference BERI systems such as physical memory maps.

Chapter 3 describes the BERI Programmable Interrupt Controller, an integrated device that supports interrupts from more peripherals than processor interrupt lines and also for interprocessor interrupts (IPIs).

Chapter 4 provides more detailed information about the implementation of the BERI prototype, including a high-level guide to its source code.

Chapter 5 documents how to check out the BERI source code, build the BERI simulator, and run the BERI unit test suite. Various build options are discussed, including debug options.

Chapter 6 describes the BERI hardware debug unit, which allows low-level access to processor internals via a real or simulated UART, for the purposes of debugging.

Chapter 7 documents the BERI unit test suite, including how to run the suite and add new tests.

Chapter 8 describes how to configure and synthesize BERI in the Altera development environment.

Chapter 9 describes how to build and synthesize the BERI prototype for the Terasic tPad and DE4 FPGA teaching boards.

Chapter 2

The BERI ISA

This chapter describes the Instruction Set Architecture (ISA) implemented by the BERI and BERI2 processors. The core CPU features are described, MIPS and CHERI ISA status are enumerated, BERI's modifications to the TLB interface, and features such as multi-threading are described.

2.1 BERI CPU Features and ISA

The intent of the BERI prototype is to support exploration and validation of the CHERI fine-grained in-address-space memory protection and scalable compartmentalization models. Our goal was not to create a complete and productizable processor design – the marketplace has many high-quality commercial embedded RISC processors. Instead we hope to provide a flexible and extensible platform for research into the hardware-software interface to facilitate the development of new ideas in processor design. To this end, BERI is prototyped in the high-level Bluespec System Verilog (BSV) Hardware Description Language (HDL), which supports highly parameterizable designs and a software-style development process.

While designing a new ISA entirely from scratch would have been possible, we instead selected a 1994-vintage version of the 64-bit MIPS ISA as a starting point which allows us to incrementally deploy and evaluate new ISA features against a known baseline, as well as demonstrate the realism of our approach. We are able to exploit extensive existing software infrastructure including compilers, toolchain, debuggers, operating systems, and applications: we are able to run the open-source FreeBSD operating system and many higher-level applications. We have implemented CHERI's capability features using the MIPS co-processor 2 instruction encoding space, and adapted FreeBSD, Clang/LLVM, and several applications to use its features.

The BERI prototype implements a set of high-level hardware features comparable to those found in the MIPS R4000 processor:

- A pipelined processor design.
- 32 64-bit general-purpose registers usable with the MIPS n64 ABI
- A full range of branch and control operations, including conditional branches, conditional traps, jump-and-link, and system calls, as well as a branch predictor.
- 16K L1 instruction and data caches which are direct-mapped, write-through, virtually indexed, and physically tagged with 32-byte lines

- 64K shared L2 cache which is direct-mapped, write-back, physically indexed and tagged with 32-byte lines.
- 64-bit integer ALU, including support for multiply and divide.
- Coprocessor 0 (CP0), system control features such as a sufficiently capable MMU to implement OS virtual memory and paging features.
- An IEEE-compatible floating-point unit (FPU)
- Multiple CPU protection rings (kernel, supervisor, usermode).
- Mature exception handling, including cycle timer, various arithmetic and memory access exceptions, and interrupt delivery from external devices.
- Programmable interrupt controller (PIC) able to multiplex larger number of interrupt sources to the smaller number of IRQ lines supported by the MIPS ISA. This will also provide future support for interprocessor interrupts (IPIs) required for multi-processor operation.

At this time, the BERI prototype omits a number of features found in the MIPS R4000, largely because they are not required for validation of the research hypotheses we are exploring. Some other modifications were made due to the specific implementation characteristics of FPGA soft cores. In particular, the decision to implement smaller caches was motivated by the performance trade-offs in the FPGA substrate which provides comparatively high-speed main memory, as well as a desire for simplicity. The following features are omitted from the MIPS 4000 ISA, or significantly modified:

- Only 64-bit addressing mode; no 32-bit addressing support.
- Only big endian support; no variable-endian features.
- Currently, BERI is a single-core, single-threaded processor; we have in-progress prototypes adding multithreading and multiprocessing support.

The following sections provide more detailed information on the 64-bit MIPS subset implemented in the BERI prototype.

2.1.1 MIPS Instructions

BERI implements roughly the instruction set found in the MIPS R4000, subject to high-level variations described in the previous section. The following tables document in greater detail the MIPS ISA instructions implemented in the BERI prototype, followed by notes on any limitations to specific instructions:

Table	Description
Table 2.1	MIPS load and store instructions
Table 2.2	MIPS arithmetic instructions
Table 2.3	MIPS logical and bitwise instructions
Table 2.4	MIPS jump and branch instructions
Table 2.5	MIPS coprocessor instructions
Table 2.6	MIPS special instructions
Table 2.7	Additional instructions from other MIPS ISA versions

Instruction	Description	Status
LB	Load byte	Implemented
LBU	Load byte unsigned	Implemented
LD	Load doubleword	Implemented
LDL	Load doubleword left	Implemented
LDR	Load doubleword right	Implemented
LH	Load halfword	Implemented
LHU	Load halfword unsigned	Implemented
LL	Load linked	Implemented
LLD	Load linked doubleword	Implemented
LUI	Load upper immediate	Implemented
LW	Load word	Implemented
LWL	Load word left	Implemented
LWR	Load word right	Implemented
LWU	Load word unsigned	Implemented
SB	Store byte	Implemented
SC	Store conditional	Implemented
SCD	Store conditional doubleword	Implemented
SD	Store doubleword	Implemented
SDL	Store doubleword left	Implemented
SDR	Store doubleword right	Implemented
SH	Store halfword	Implemented
SW	Store word	Implemented
SWL	Store word left	Implemented
SWR	Store word right	Implemented

Table 2.1: MIPS load and store instructions in the BERI prototype

Instruction	Description	Status
ADD	Add	Implemented
ADDI	Add immediate	Implemented
ADDIU	Add immediate unsigned	Implemented
ADDU	Add unsigned	Implemented
ANDI	And immediate	Implemented
DADD	Doubleword add	Implemented
DADDI	Doubleword immediate	Implemented
DADDIU	Doubleword add immediate unsigned	Implemented
DADDU	Doubleword add unsigned	Implemented
DDIV	Doubleword divide	Implemented
DDIVU	Doubleword divide unsigned	Implemented
DIV	Divide	Implemented
DIVU	Divide unsigned	Implemented
DMULT	Doubleword multiple	Implemented
DMULTU	Doubleword multiple unsigned	Implemented
DSUB	Doubleword subtract	Implemented
DSUBU	Doubleword subtract unsigned	Implemented
MFHI	Move from HI	Implemented
MFLO	Move from LO	Implemented
MTHI	Move to HI	Implemented
MTLO	Move to LO	Implemented
MULT	Multiply	Implemented
MULTU	Multiply unsigned	Implemented
SLT	Set on less than	Implemented
SLTI	Set on less than immediate	Implemented
SLTIU	Set on less than immediate unsigned	Implemented
SLTU	Set on less than unsigned	Implemented
SUB	Subtract	Implemented
SUBU	Subtract unsigned	implemented

Table 2.2: MIPS arithmetic instructions in the BERI ISA

Instruction	Description	Status
AND	And	Implemented
DSLL	Doubleword shift left logical	Implemented
DSLLV	Doubleword shift left logical variable	Implemented
DSLL32	Doubleword shift left logical + 32	Implemented
DSRA	Doubleword shift right arithmetic	Implemented
DSRAV	Doubleword shift right arithmetic variable	Implemented
DSRA32	Doubleword shift right arithmetic + 32	Implemented
DSRL	Doubleword shift right logical	Implemented
DSRLV	Doubleword shift right logical variable	Implemented
DSRL32	Doubleword shift right logical + 32	Implemented
NOR	Nor	Implemented
OR	Or	Implemented
ORI	Or immediate	Implemented
SLL	Shift left logical	Implemented
SLLV	Shift left logical variable	Implemented
SRA	Shift right arithmetic	Implemented
SRAV	Shift right arithmetic variable	Implemented
SRL	Shift right logical	Implemented
SRLV	Shift right logical variable	Implemented
XOR	Exclusive or	Implemented
XORI	Exclusive or immediate	Implemented

Table 2.3: MIPS logical and bitwise instructions in the BERI ISA

Instruction	Description	Status
BEQ	Branch on equal	Implemented
BEQL	Branch on equal Likely	Implemented
BGEZ	Branch on greater than or equal to zero	Implemented
BGEZAL	Branch on greater than or equal to zero and link	Implemented
BGEZALL	Branch on greater than or equal to zero and link likely	Implemented
BGEZL	Branch on greater than or equal to zero likely	Implemented
BGTZ	Branch on greater than zero	Implemented
BGTZL	Branch on greater than zero likely	Implemented
BLEZ	Branch on less than or equal to zero	Implemented
BLEZL	Branch on less than or equal to zero likely	Implemented
BLTZ	Branch on less than zero	Implemented
BLTZAL	Branch on less than zero and link	Implemented
BLTZALL	Branch on less than zero and link likely	Implemented
BLTZL	Branch on less than zero likely	Implemented
BNE	Branch on not equal	Implemented
BNEL	Branch on not equal likely	Implemented
J	Jump	Implemented
JAL	Jump and link	Implemented
JALR	Jump and link register	Implemented
JR	Jump register	Implemented

Table 2.4: MIPS jump and branch instructions in the BERI ISA

Instruction	Description	Status
BCzF	Branch on Coprocessor z False	Not implemented
BCzFL	Branch on Coprocessor z False Likely	Not implemented
BCzT	Branch On Coprocessor z True	Not implemented
BCzTL	Branch On Coprocessor z True Likely	Not implemented
CFCz	Move control from coprocessor	Not implemented
COPz	Coprocessor operation	See Section 2.1.4
CTCz	Move control to coprocessor	Not implemented
DMFC0	Doubleword move from system control coprocessor	See Section 2.1.2
DMTC0	Doubleword move to system control coprocessor	See Section 2.1.2
LDCz	Load doubleword to coprocessor	Not implemented
LWCz	Load word to coprocessor	Not implemented
MFC0	Move from system control coprocessor	See Section 2.1.2
MFCz	Move from coprocessor	See Section 2.1.4
MTC0	Move to system control coprocessor	See Section 2.1.2
MTCz	Move to coprocessor	See Section 2.1.4
SDCz	Store doubleword from coprocessor	Not implemented
SWCz	Store word from coprocessor	Not implemented
TLBP	Probe TLB for matching entry	Implemented
TLBR	Read indexed TLB entry	Implemented
TLBWI	Write indexed TLB entry	Implemented
TLBWR	Write random TLB entry	Implemented

Table 2.5: MIPS coprocessor instructions in the BERI ISA. See Section 2.1.2 for limitations on CP0 instructions; see Section 2.1.4 for limitations on generic coprocessor instructions

Instruction	Description	Status
BREAK	Breakpoint	Implemented
CACHE	Cache	Implemented
ERET	Exception return	Implemented
SYNC	Synchronize	See Section 2.1.6
SYSCALL	System call	Implemented
TEQ	Trap if equal	Implemented
TEQI	Trap if equal immediate	Implemented
TGE	Trap if greater than or equal	Implemented
TGEI	Trap if greater than or equal immediate	Implemented
TGEIU	Trap if greater than or equal immediate unsigned	Implemented
TGEU	Trap if greater than or equal unsigned	Implemented
TLT	Trap if less than	Implemented
TLTI	Trap if less than immediate	Implemented
TLTIU	Trap if less than immediate unsigned	Implemented
TLTU	Trap if less than unsigned	Implemented
TNE	Trap if not equal	Implemented
TNEI	Trap if not equal immediate	Implemented

Table 2.6: MIPS special instructions in the BERI ISA

Instruction	Description	Version	Status
MADD	Multiply and add signed words to HI, LO	MIPS32	Implemented
MADDU	Multiply and add unsigned words to HI, LO	MIPS32	Implemented
MOVN	Move conditional on not zero	MIPS32	Implemented
MOVZ	Move conditional on zero	MIPS32	Implemented
MSUB	Multiply and subtract from HI, LO	MIPS32	Implemented
MSUBU	Multiply and subtract from HI, LO	MIPS32	Implemented
MUL	Multiply word to general-purpose register	MIPS32	<i>Untested</i>
RDHWR	Read hardware register	MIPS32R2	Implemented
SSNOP	Superscalar no operation	MIPS32	See Section 2.1.6
WAIT	Enter standby mode	MIPS32	Not implemented

Table 2.7: Selected additional instructions in the BERI ISA, derived from other MIPS ISA versions; see Section 2.1.6

Rounding mode	Status
Round to nearest, tie to even	Implemented
Round towards zero	Not implemented
Round towards $+\infty$	Not implemented
Round towards $-\infty$	Not implemented

Table 2.8: IEEE floating point rounding modes supported by the BERI ISA

Type	Status
Single precision	Implemented
Double precision	Implemented
Paired single	Implemented

Table 2.9: IEEE floating point types supported by the BERI ISA

Instruction	Status
ADD	Implemented
SUB	Implemented
MUL	Implemented
DIV	Implemented
ABS	Implemented
MOV	Implemented
NEG	Implemented
SQRT	Implemented

Table 2.10: FPU computational instructions supported by the BERI ISA

Instruction	Status
CVT.S.D	Implemented
CVT.S.W	Implemented
CVT.S.L	Implemented
CVT.D.S	Implemented
CVT.D.W	Implemented
CVT.D.L	Implemented
CVT.W.S	Implemented
CVT.W.D	Implemented
CVT.L.S	Implemented
CVT.L.D	Implemented
ROUND.W.S	Implemented
ROUND.W.D	Implemented
ROUND.L.S	Implemented
ROUND.L.D	Implemented
TRUNC.W.S	Implemented
TRUNC.W.D	Implemented
TRUNC.L.S	Implemented
TRUNC.L.D	Implemented
CEIL.W.S	Implemented
CEIL.W.D	Implemented
CEIL.L.S	Implemented
CEIL.L.D	Implemented
FLOOR.W.S	Implemented
FLOOR.W.D	Implemented
FLOOR.L.S	Implemented
FLOOR.L.D	Implemented

Table 2.11: Floating point conversion instructions supported by the BERI ISA

Instruction	Version	Status
LDXC1	MIPS IV	Implemented
LWXC1	MIPS IV	Implemented
MOVF	MIPS IV	Implemented
MOVF.D	MIPS IV	Implemented
MOVF.S	MIPS IV	Implemented
MOVN.D	MIPS IV	Implemented
MOVN.S	MIPS IV	Implemented
MOVT	MIPS IV	Implemented
MOVZ.D	MIPS IV	Implemented
MOVZ.S	MIPS IV	Implemented
RECIP.D	MIPS IV	Implemented
RECIP.S	MIPS IV	Implemented
RSQRT.D	MIPS IV	Implemented
RSQRT.S	MIPS IV	Implemented
SDXC1	MIPS IV	Implemented
SWXC1	MIPS IV	Implemented

Table 2.12: Floating point instructions from later MIPS versions supported by BERI

2.1.2 Limitations to Coprocessor 0 Support

Currently, not all system control coprocessor (CP0) features are implemented in the BERI prototype. In the future, this section will explicitly document included CP0 features.

2.1.3 Modifications to the MIPS TLB Model

The MIPS R4000 MMU implements a 48-entry, fully associative Translation Look-aside Buffer (TLB). Software interacts with the MIPS R4000 TLB by performing Write Indexed or Write Random operations; the latter operations use the Random CP0 register contents as the index. The Random CP0 register decrements every cycle but resets to the highest TLB entry when it reaches the Wired CP0 register. Thus, a Write Random operation never overwrites TLB entries below the Wired register.

BERI's TLB is configurable but is currently composed of a lower-16 group of fully associative entries and an upper-128 group of direct-mapped entries. We use this configuration because, whereas we desire a large TLB, FPGA fabrics are unable to efficiently construct large associative searches. This associative plus mapped structure allows an arbitrarily large TLB with trivial additional logic, because most of the entries are stored in block RAM.

Our implementation behaves differently from a simple MIPS R4000 TLB in several ways:

1. Writes of arbitrary indexed values are supported for only the lower 16 entries. An indexed write to the upper 128 entries will result in a write to an index above 15 whose lower 7 bits are equal to the lower 7 bits of the virtual page number.

2. Write Random operations will not write to a random location but rather to an index above 15 whose lower 7 bits are equal to the lower 7 bits of the virtual page number of the written entry.
3. A valid entry that is displaced by a Write Random instruction will be placed in an unpredictable location above the wired entry and less than or equal to 15. Thus, the fully associative entries that are not wired act as a victim buffer for the direct-mapped entries.
4. BERI only supports variable sized pages in the lower associative entries.
5. BERI's TLB implementation also does not support 32-bit virtual addresses as MIPS R4000 does.
6. BERI's TLB supports 40-bit physical addresses instead of 36-bit physical addresses in MIPS R4000. This means the EntryLo registers have a 28-bit PFN field, and the size of the EntryLo registers is 34 bits in total.

BERI's TLB implementation works for FreeBSD without modification. FreeBSD wires only the bottom TLB entry for use in exception handlers, and then accesses the rest of the TLB entries chiefly using the Write Random operation – with Write Indexed operations being used only to modify entries in place or to invalidate TLB entries. When FreeBSD invalidates TLB entries, it uses a virtual page number whose lower bits are equal to the index number in the TLB. Thus, our design (which takes the lower bits of the virtual page number as the index) works as expected. Furthermore, FreeBSD always probes to find the index of an entry immediately before modifying it, and does not remember where it placed entries in the table. Thus, FreeBSD is not confused when our implementation relocates entries from the direct-mapped region to the victim buffer.

If any operating system or hypervisor desired to more closely manage the TLB, it should take into account the mapped nature of the upper entries of the TLB and the possibility that a Write Random operation may relocate a previously mapped entry.

In other respects, the BERI TLB is similar to the MIPS R4000, including the MIPS design choice in which each TLB entry maps two pages.

2.1.4 Limitations to Generic Coprocessor Support

The BERI prototype does not currently support generic coprocessor instructions, but does implement an interface for the capability coprocessor using coprocessor 2.

2.1.5 The BERI Floating-Point Unit (FPU)

The BERI floating point unit is a fairly complete implementation of the MIPS R4000 floating point instructions, with the following omissions:

- The only supported rounding mode is round to nearest even.
- Floating point exceptions are not implemented. Where appropriate, instructions will return an IEEE “infinity” or “not a number” value, but it is not possible to enable exceptions for these cases.
- Floating point status flags are not implemented.

The FPU also implements some instructions from MIPS IV; see table 2.12. Many of these instructions are used by the Clang/LLVM compiler.

In conformance with the MIPS R4000 specification, the `abs.s`, `abs.d`, `neg.s` and `neg.d` instructions are what the IEEE 754-1985 standard calls “arithmetic” operations. Later revisions of the MIPS specification introduced a control bit in FCSR, ABS2008, which causes `abs` and `neg` to be “non-arithmetic” instructions as required by IEEE 754-2008. We do not implement the ABS2008 control bit. Because they are “arithmetic”, computing the `abs` or `neg` of a signalling NaN would raise an invalid operation exception if floating point exceptions were enabled. (But we do not support enabling this exception, as described above). For the same reason, `abs` or `neg` of a quiet NaN does not change the sign bit.

2.1.6 Selected Additions from Later MIPS ISA Versions

Table 2.7 documents selected instructions added to the BERI ISA from later MIPS versions. In general, we have added instructions only where required by common compiler toolchains and operating systems.

In that the BERI prototype is pipelined, but currently neither superscalar nor multicore, the `SSNOP` and `SYNC` instructions are interpreted as a NOPs. If and when superscalar support is added to future BERI versions, that support will need to be enhanced.

The BERI prototype implements the `config1` CP0 shadow register, introduced in MIPS32, which allows queries of cache layout properties. This is used by FreeBSD during CPU discovery to select cache management routines.

The BERI prototype also implements the `RDHWR` (read hardware register) instruction. The registers which can be read using this instruction are the CP0 count register and the “user local” register. The `ULRI` bit of `config` register 3 is set to indicate that the user local register is implemented. The user local register can be written as CP0 register 4, select 2.

The CP0 register `config6` is left as implementation-defined in the MIPS specification. In BERI, it is defined as follows:

Bits	Description
16–31	Size of TLB-1
3–15	Zero
2	Enable large TLB
0–1	Zero

2.1.7 Virtual Address Space

The 64-bit MIPS ISA divides its 64-bit address space into a number of segments with various properties. BERI implements roughly the same address space layout as the MIPS R4000, except that CP0 status register bits **UX**, **SX**, and **KX** are always set to 1. This means that user mode, supervisor mode, and kernel mode must always execute in 64-bit addressing mode in BERI (i.e., no 32-bit addressing is supported).

Ring 2: User Mode

Table 2.13 illustrates the user mode address space. The processor is in user mode whenever **KSU** is 10, **EXL** is 0, and **ERL** is 0. In this mode, only the lower quarter of the address space is available and all addresses are virtual and mapped by the TLB.

Start address	Stop address	Description	Size
0x40000000 00000000	0xffffffff ffffffff	address error	
0x00000000 00000000	0x3fffffffff ffffffff	xuseg (user)	2 ⁴⁰ Bytes

Table 2.13: BERI address space layout in user mode

Start address	Stop address	Description	Size
0xffffffff e0000000	0xffffffff ffffffff	address error	
0xffffffff c0000000	0xffffffff dfffffff	csseg	512M
0x40000000 00000000	0x400000ff ffffffff	xsseg	2 ⁴⁰ Bytes
0x00000000 00000000	0x000000ff ffffffff	xsuseg (user)	2 ⁴⁰ Bytes

Table 2.14: BERI address space layout in supervisor mode

Ring 1: Supervisor Mode

Table 2.14 illustrates the supervisor address space. The processor is in supervisor mode whenever **KSU** is 01, **EXL** is 0, and **ERL** is 0. All available addresses are virtual and mapped by the TLB. Unavailable addresses give an address error exception when referenced.

Ring 0: Kernel Mode

Table 2.15 illustrates the kernel address space. Table 2.16 details the `xkphys` subset of the address space, in which the physical memory space is mapped directly into regions of virtual address space – using various caching policies. The processor is in kernel mode if **KSU** is 0, **EXL** is 1, or **ERL** is 1.

Start address	Stop address	Description	Size
0xffffffff e0000000	0xffffffff ffffffff	ckseg3 - mapped	512M
0xffffffff c0000000	0xffffffff dfffffff	cksseg - mapped	512M
0xffffffff a0000000	0xffffffff bfffffff	ckseg1 - unmapped, uncached	512M
0xffffffff 80000000	0xffffffff 9fffffff	ckseg0 - unmapped, cached	512M
0xc0000000 00000000	0xc00000ff 7fffffff	xkseg - mapped	
0x80000000 00000000	0x800000ff ffffffff	xkphys - unmapped	
0x40000000 00000000	0x400000ff ffffffff	xsseg	2 ⁶² Bytes
0x00000000 00000000	0x000000ff ffffffff	xsuseg (user)	2 ⁶² Bytes

Table 2.15: BERI address space layout in kernel mode

Start address		Stop address		Description
0xb8000000	00000000	0xb80000ff	ffffffff	reserved
0xb0000000	00000000	0xb00000ff	ffffffff	cached, coherent update on write
0xa8000000	00000000	0xa80000ff	ffffffff	cached, coherent exclusive on write
0xa0000000	00000000	0xa00000ff	ffffffff	cached, coherent exclusive
0x98000000	00000000	0x980000ff	ffffffff	cached, noncoherent
0x90000000	00000000	0x900000ff	ffffffff	uncached
0x88000000	00000000	0x880000ff	ffffffff	reserved
0x80000000	00000000	0x800000ff	ffffffff	reserved

Table 2.16: Layout of the xkphys region in BERI

Instruction	Description	Status
CGetBase	Move base to a general-purpose register	Implemented
CGetLen	Move length to a general-purpose register	Implemented
CGetPerm	Move permissions field to a general-purpose register	Implemented
CGetType	Move object type field to a general-purpose register	Implemented
CGetUnsealed	Move unsealed flag to a general-purpose register	Implemented
CGetPCC	Move the PCC and PC to a general-purpose registers	Implemented
CGetTag	Move the valid capability flag to a general-purpose register	Implemented
CSetType	Set the object type field of an executable capability	Implemented
CIncBase	Increase Base	Implemented
CSetLen	Decrease Length	Implemented
CAndPerm	Restrict Permissions	Implemented
CClearTag	Clear the capability valid flag	Implemented

Table 2.17: CHERI ISA instructions for getting and setting capability register fields

2.2 CHERI ISA Extensions

The BERI and BERI2 prototypes have undergone a number of revisions as the CHERI ISA matured, and now fully implement the feature set described in the *CHERI Architecture Document*, including capability co-processor instructions, exception model, and tagged memory. CHERI Clang/LLVM and CheriBSD are compiled to use these features, and able to demonstrate the ISA’s support for both memory protection and sandboxing.

The following tables document in greater detail the CHERI ISA instructions implemented in the BERI prototype, followed by notes any limitations to specific instructions:

Table	Description
Table 2.17	Getting and setting capability fields
Table 2.18	Loading and storing [via] capabilities
Table 2.19	Instructions relating to object capabilities

2.3 Other BERI ISA Extensions

Apart from the CHERI capability instructions (described in the previous section) BERI makes the following extensions to the MIPS ISA:

- CP0 register 8, selector 1: After an exception, this CP0 register will contain the instruction that caused the exception.

Instruction	Description	Status
CSC	Store Capability	Implemented
CLC	Load Capability	Implemented
CLB	Load Byte Via Capability Register	Implemented
CLH	Load Half-Word Via Capability Register	Implemented
CLW	Load Word Via Capability Register	Implemented
CLD	Load Double Via Capability Register	Implemented
CLBU	Load Byte Unsigned via Capability Register	Implemented
CLHU	Load Half-Word Unsigned via Capability Register	Implemented
CLWU	Load Word Unsigned via Capability Register	Implemented
CSB	Store Byte Via Capability Register	Implemented
CSH	Store Half-Word Via Capability Register	Implemented
CSW	Store Word Via Capability Register	Implemented
CSD	Store Double Via Capability Register	Implemented
CLLD	Load Linked Double Via Capability Register	Implemented
CSCD	Store Conditional Double Via Capability Register	Implemented

Table 2.18: CHERI ISA instructions for loading and storing [via] capabilities

Instruction	Description	Status
CGetCause	Read capability exception cause register	Implemented
CSetCause	Set capability exception cause register	Implemented
CJR	Jump Capability Register	Implemented
CJALR	Jump and link Capability Register	Implemented
CBTS	Branch if tag bit is set	Implemented
CBTU	Branch if tag bit is not set	Implemented
CSealCode	Seal an executable capability	Implemented
CSealData	Seal a non-executable capability with the object type of an executable capability	Implemented
CUnseal	Unseal a sealed capability	Implemented
CCall	Protected procedure call into a new security domain	Implemented
CReturn	Return to the previous security domain	Implemented
CGetCause	Return the capability-cause register	Implemented
CSetCause	Set the capability-cause register	Implemented
CCheckPerm	Check capability permissions	Implemented
CCheckType	Check capability type	Implemented

Table 2.19: CHERI ISA instructions for creating and invoking object capabilities

2.4 BERI2 Specific Features

In general BERI2 implements the same MIPS subset and capability extensions as documented in the previous sections. However, it also implements some other experimental features, and there are some minor differences as described here. (Note that BERI2 is still a rapidly developing prototype. As a consequence, it is wise to check with the developers before relying on any of this information.)

2.4.1 Multi-Threading

BERI2 implements a simple form of multi-threading. It supports a configurable, statically determined number of hardware threads which are rotated on a cycle-by-cycle basis (barrel scheduled). In the current implementation all threads are runnable at all times, so there is no way to suspend a thread in hardware.

To build with multiple threads, edit the `ThreadSZ` definition in `cheri2/trunk/CHERITypes.bsv`. This definition sets the number of bits of thread ID, which determines the number of threads, $T = 2^{\text{ThreadSZ}}$. By default, this is set to 0 for single-threaded execution.

Most supported CP0 registers are maintained independently for each thread, with the following notable exceptions:

PrID (15) Read-only, bits 31:24 contain current thread ID.

COUNT (9) Read-only and consistent between threads. Each thread has its own COMPARE register.

CAUSE:IP6.2 (13) External hardware interrupts are delivered to all threads; however, the built-in timer interrupt (IP7) and software interrupt bits (IP0 and IP1) are maintained for each thread. Each thread can control the delivery of interrupts individually via the mask field in the status register.

TLB The TLB is currently shared between all threads. Although the relevant CP0 registers are replicated per thread, the software must implement synchronization around TLB operations in order to avoid race conditions (e.g., between `PROBE` and `TLBR` or `TLBWI` instructions).

Chapter 3

The BERI Programmable Interrupt Controller

This chapter describes the Programmable Interrupt Controller (PIC) attached to each BERI core. The PIC provides a simple way for mapping a potentially large number of external interrupts onto the small set of hardware interrupts defined by the MIPS ISA (see Table 3.3). On BERI2 each interrupt must also be mapped to a particular hardware thread. The PIC exposes memory mapped registers which can be used by software to configure the mapping and also to set, clear and read pending interrupts. Thus, the PIC allows interrupts to be triggered by both hardware-wired peripherals (e.g., a UART) and by software, referred to respectively as hard and soft sources. This latter facility can be used for inter-processor interrupts (IPIs) on multi-threaded and multicore configurations.

3.1 Sources

The PIC consists of S sources, which may be either hard or soft.

Soft The value of a soft source comes from its interrupt pending (IP) state bit, which can be set or cleared by software. In the future, some external event, such as a message received over the inter-core interconnect may potentially set these bits.

Hard The value of a hard source comes directly from a peripheral device, and is not latched. Hard sources also have an IP bit which may be manipulated in the same way as for soft sources, mainly for debugging purposes.

To calculate the current state for a particular MIPS interrupt, the PIC ORs the value of all sources which are enabled and mapped to that interrupt.

3.2 Source Numbers and Base Addresses on BERI

The source numbers and register base addresses are as shown in Tables 3.1 and 3.2.

Source No	Use
0–63	Up to 64, hard-wired external interrupts
64–127	64 implemented soft interrupts
128–1023	Reserved soft interrupts, unimplemented

Table 3.1: BERI PIC source number allocation

Register Name	Address	Used to
PIC_CONFIG_BASE	Refer to relevant section of Chapter 9	Configure interrupts source mappings
PIC_IP_READ_BASE	PIC_CONFIG_BASE + 8 * 1024	Read interrupt source state
PIC_IP_SET_BASE	PIC_IP_READ_BASE + 128	Set interrupt pending bits
PIC_IP_CLEAR_BASE	PIC_IP_READ_BASE + 256	Clear interrupts pendings btis

Table 3.2: BERI PIC control register addresses

3.3 Config Registers: PIC_CONFIG_X

Each source has an associated configuration register with the format shown in Figure 3.3. This register allows direction of the interrupt to a given interrupt number of a given thread; it also can enable or disable delivery of the interrupt. The configuration register for source s has address $\text{PIC_CONFIG_BASE} + 8s$. The default value for all configuration fields is 0 (i.e. disabled). Word or double-word accesses may be used.

3.4 Interrupt Pending Bits

Each interrupt source has one associated bit for an interrupt pending (IP) state. For hardware sources, the IP bit is ORed with the incoming interrupt wire to provide the current value for the interrupt source. This bit may be used to artificially trigger an interrupt for debugging purposes. Note that the hardware interrupt is *not* latched by the IP bit; therefore, the source will stay high only as long as the hardware source asserts its interrupt, and will go low once software has dealt with the interrupt at the device. This behavior is consistent with the IP bits in the MIPS cause register.

The PIC also provides soft sources which may be used for inter-thread interrupts. We expect that software will configure at least one soft source per thread for this use. If non-maskable or debug inter-thread interrupts are also required then two or more sources per thread may be configured.

On future multi-processor builds, a message on the inter-processor interconnect will be able to set an interrupt pending bit to be set, thus allowing for inter-processor interrupts or message-based interrupts similar to PCI's Message Signalled Interrupts.

The IP bits are packed into 64-bit registers for manipulation by software. The current value for a source, s , can be read from a read-only register at address $\text{PIC_IP_READ_BASE} + \lfloor s/8 \rfloor$, in bit $s \bmod 64$ (numbered from zero as the least significant bit). For hard sources, this value is the value of the external interrupt wire ORed together with the IP bit. Thus, the state of the IP bit cannot be read in isolation. Software may set the IP bit for a source by writing a value of one to the corresponding offset from PIC_IP_SET_BASE and,

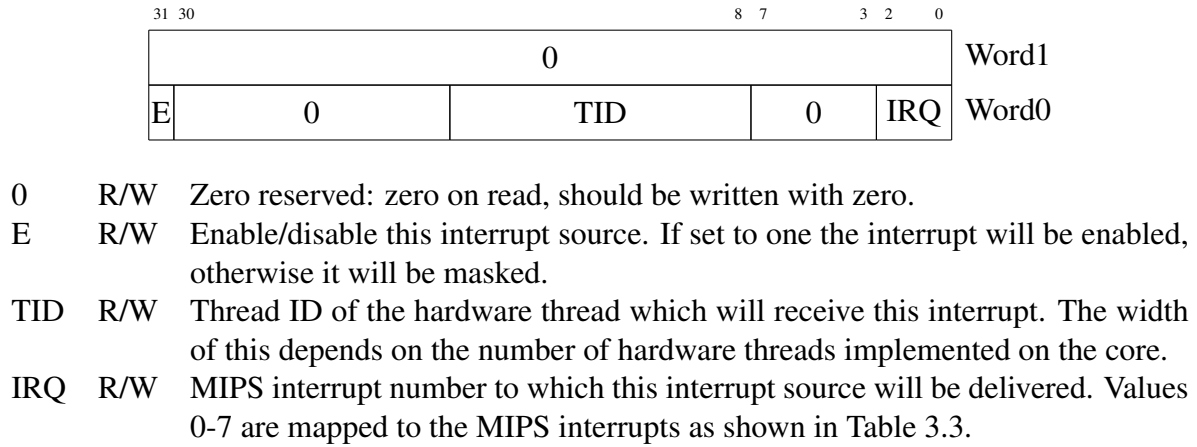


Figure 3.1: PIC Configuration Register Format

IRQ Field Value	MIPS Interrupt
0-4	MIPS external interrupts 0-4 corresponding to IP2-IP6 in the CP0 Cause register.
5	MIPS external interrupt 5 corresponding to IP7 in the CP0 Cause register. The value of IP7 is the logical OR of the output of the PIC with the core's timer interrupt value.
6	MIPS non-maskable interrupt corresponding to NMI field of the CP0 Status register. CURRENTLY NOT SUPPORTED ON BERI OR BERI2.
7	MIPS EJTAG debug exception trigger. CURRENTLY NOT SUPPORTED ON BERI OR BERI2.

Table 3.3: Values of the IRQ field of config register.

similarly, clear it using an offset from `PIC_IP_CLEAR_BASE`. Bits written with zero will have no effect. The set/clear semantics allows for atomic manipulation of one or more bits in the packed registers without the potential for race conditions associated with a read/modify/write sequence.

3.5 Reset State

On reset, all PIC configuration and state is set to zero except for the first five hardware sources, 0-4, which are given a backwards-compatible initial state as follows:

- The E bit is set to one to enable the interrupt
- The TID field is set to zero to pass the interrupt to thread 0
- The IRQ field is set to the source number

Effectively, the PIC is completely transparent to PIC-unaware code, which may behave as if external interrupts were directly connected to the MIPS interrupt wires. PIC aware software should not rely on this behavior and should explicitly configure all interrupt sources on boot.

3.6 Safe Handling of Interrupts

A combination of soft interrupts and shared-memory communication to pass inter-thread or inter-processor messages is likely to be used. In order to do this safely while avoiding missed wakeups, the source should be cleared first, before handling any incoming messages in a loop. Otherwise, a spurious interrupt could result in the case where a second interrupt arrives during the processing of the first interrupt, although that would not result in missed wakeups.

For hardware sources, the device must provide a safe way of handling and quiescing the interrupt.

Chapter 4

The BERI Processor Implementation

This chapter provides a high-level overview of the BERI prototype processor implementation, including programming language choice, directory layout, and a high-level description of the prototype’s files and modules.

4.1 Bluespec

The BERI prototype is implemented in the Bluespec Hardware Description Language (HDL), a Haskell-derived programming language that allows highly parameterized and structured logic designs. Bluespec source code may be compiled to an efficient C simulation, or into Verilog for simulation or synthesis. One of the key properties of Bluespec is that it makes design space exploration far more accessible than traditional low-level HDLs, which is critical for fast and easy evaluation of the impact of our design choices on a practical hardware implementation.

4.2 Directory Layout

The Bluespec source code for the BERI processor resides in the root of the BERI distribution. A series of sub-directories, listed in Table 4.1, contain a combination of supplementary software source code including an interactive self-test and unit test suite and tools used in building BERI. These subdirectories are also targets for generated files.

Directory	Description
<code>trunk/</code>	Root of the BERI source tree, home of Bluespec source code
<code>boards/</code>	Holds project directories for various FPGA boards
<code>sw/</code>	Home of integrated software component source code
<code>ip/</code>	Destination for generated Verilog files

Table 4.1: Directories in the BERI source code distribution

4.3 Key Files

Table 4.2 describes the key files in the BERI prototype implementation.

4.4 Conditionally Compiled Features

Table 4.3 is a list of conditionally compiled features in the BERI processor. Values for these macros are selected by various build targets described in Chapter 5, and passed to the Bluespec compiler via the `bsc` command line.

File	Description
Makefile	GNU makefile to build BERI
MIPS.bsv	Types and shared functions for the design
MIPSTop.bsv	Top-level module implementing instruction and register fetch, which instantiates all other modules
Scheduler.bsv	Pre-decode stage of the pipeline
Decode.bsv	Decode stage of the pipeline
Execute.bsv	Execute stage of the pipeline
MemAccess.bsv	Memory access and writeback stages of the pipeline
Memory.bsv	Memory subsystem, which instantiates the caches, merging logic, and memory interface
ICache.bsv	Instruction level 1 cache
DCache.bsv	Data level 1 cache
Merge.bsv	Module merging memory interfaces into a single interface in the memory heirarchy
L2Cache.bsv	L2 cache
TopAvalonPhy.bsv	Top-level module adapting BERI's memory interface to an Avalon bus interface
TopSimulation.bsv	Top-level module interfacing BERI's memory interface with the PISM bus for C peripheral models
MIPSRegFile.bsv	Forwarding register file
CP0.bsv	Coprocessor 0 containing all configuration registers
TLB.bsv	40-entry TLB with three cached interfaces
CapCop.bsv	Module implementing the capability coprocessor

Table 4.2: Key files in the BERI source code

Macro	Description
CAP	Include capability coprocessor
COP1	Include floating point unit
MULTI	Enable multi-core
MULTICORE	Number of cores
MICRO	Do not include the TLB and L2 Cache

Table 4.3: Macros controlling conditionally compiled features in BERI1

Macro	Description
CAP	Include capability coprocessor
DEBUG	Enable hardware-level tracing
THREADSZ	log2 of the number of threads
VERIFY	Build slower, but easier to verify, version

Table 4.4: Macros controlling conditionally compiled features in BERI2

Chapter 5

Simulating BERI

This chapter describes how to check out the BERI source code, build the BERI simulator, and run the BERI unit test suite. It describes various build targets for the simulator with varying levels of tracing. This documentation assumes access to the following resources:

- 32-bit or 64-bit Ubuntu Linux 10.04.2 LTS or 12.04 LTS workstation or server
- Access to the hosted Subversion repository

5.1 Software Dependencies to Build BERI

Table 5.1 documents software packages required to build the BERI simulator and test suite. For Altera tools to function under Ubuntu, `/bin/sh` must point to `bash`, and not `dash`, which is the default in Ubuntu. This can be accomplished with the following commands:

```
# rm /bin/sh
# ln -s /bin/sh bash
```

To install an Ubuntu package such as those listed in the table, using the following command:

```
# apt-get install <package-name>
```

If Ubuntu 10.04 (Lucid) or 9.10 (Karmic) is used, the backports source must be enabled in `/etc/apt/sources.list` to install `libstdc++5`.

The SDE cross-compiling distribution of the GCC compiler on Linux should be available from MIPS.com. The current URL is unknown.

The Cambridge extended assembler is available from `/usr/groups/ctsrld` on the network filer at the Cambridge Computer Laboratory.

Program	Ubuntu Package	Required to
Subversion	subversion	Check out
GNU make	make	Build, run test suite
GNU Bison 2.4.1	bison	Build simulator
Flex 2.5.35	flex	Build simulator
Bluespec	See Bluespec	Build
2011.05.beta1	Documentation	
GNU C++ libraries (v5)	libstdc++5	Running Bluespec
Perl	perl	Build
SDL 1.2.14	libsdl1.2-dev, libsdl1.2debian	Build, run simulated tPad frame buffer
Python 2.6	python	Build, run test suite
SDE toolchain 6.06	none	Build test suite
Cambridge extended assembler	none	Build test suite
Nose 0.11.1	python-nose	Run test suite
Ubuntu 32-bit libraries	ia32-libs	Build test suite on 64-bit systems

Table 5.1: Software build dependencies for BERI components

5.2 Checking out BERI

The Cambridge Subversion repository uses SSH authentication keys as capabilities to identify the repository and rights held by a client. By default, SSH will offer keys held by the agent (or in your home directory) in the order it finds them, which, if you hold multiple keys to different repositories on the Subversion server, may cause SSH to select the wrong key. It is therefore necessary to ensure that the right SSH key is used. One way to do so is to create a new SSH agent, adding only the appropriate key to that session¹. To set up an SSH agent in this manner, use something like the following:

```
$ ssh-agent bash
$ ssh-add ~/.ssh/id_ctsrd_rsa
```

It is also possible to configure `.ssh/config` to offer only a specific key to specific servers; see the SSH man page for details. To perform an initial checkout of BERI, use the following Subversion command:

```
$ svn co svn+ssh://secsvn@svn-ctsrd.cl.cam.ac.uk/ctsrd ctsrd
```

To update an existing checkout of BERI, use the following Subversion command:

```
$ cd ctsrd
$ svn update
```

5.3 Building the BERI Simulator

The BERI source code and build tools may be found in the `ctsrd/cheri/trunk` directory tree. Before building BERI, you must configure the Bluespec development environment:

```
$ cd ctsrd/cheri/trunk
$ source setup.sh
```

The BERI build is sensitive to a number of `make` variables, depending on what components should be included or excluded from the pipeline. A typical run of the simulator might be initiated using:

```
$ CAP=1 make sim
$ ./sim
```

The BERI build is sensitive to five `make` variables.

If `MICRO` is set, BERI will build without L2 cache or memory translation.

If `NOBRANCH` is set, BERI will not predict branches, but will always wait for the final branch destination before fetching the next PC.

If `CAP` is set, CHERI will be built, that is, the capability unit will be included in the build.

If `COP1` is set, BERI will be built with the floating point unit.

If `COP3` is set, BERI will be built with an experimental general purpose coprocessor 3.

¹On Mac OS X, new `ssh-agent` sessions inherit all SSH keys added to the user keychain, so you must run `ssh-add -D` to flush them. This step is not required on other platforms.

The BERI executable is sensitive to four arguments. The `+trace` argument will give a concise report for each instruction committed. The `+cTrace` argument will report the number of dead cycles between committed instructions. The `+regDump` argument will enable the debug instructions which report the contents of the register files. The `+debug` argument will report all debug output from the internal processor state.

The build also compiles an interactive software test tool, found in the `sw` sub-directory. Other available software packages, stored elsewhere in the repository, include the BERI unit test suite described in Chapter 7 (`cheritest/trunk`) and the BERI demonstration environment described in the *CHERI User's Guide* (`deimos/trunk`). When the simulator is run, it loads a memory image from the current working directory; running the simulator from the root of the BERI development tree will automatically load the interactive test tool.

5.4 Configuring the BERI Simulator

At startup, the simulator tries to read the file `./simconfig` or the file pointed to by the environment variable `CHERI_CONFIG`. This file describes in a C-like syntax how the hardware should be simulated. A valid configuration must exist or the simulator will not start. A default configuration file is included in the `cheri/trunk` directory.

Individual simulated hardware peripherals are built as shared libraries. The simulator will attempt to `dlopen()` these shared libraries as they are encountered in the configuration file. Any module-specific options are passed to the module at load-time. If a module fails to load, either because it cannot be found or because invalid options were given, the simulation will terminate with an assertion failure. Figure 5.1 illustrates a sample `simconfig` file.

First, a series of simulated device modules are loaded using `module` statements; paths must be specified. Then, a series of devices is declared using `device` blocks, which must each declare a `class`, which selects the simulated device type; for each device, at least a base address (`addr`) and length (`length`) must be specified. An optional `irq` can be set, as well as device class-specific parameters such as socket types, file paths, and so on. Devices can be conditionally defined based on whether environmental variables have been set; both positive `ifdef` and negative `ifndef` syntaxes are permitted. Finally, options can be set from environmental variables using the `getenv` syntax; if a variable is not set, then an empty string will be used for the value.

5.5 Simulating BERI

When the `sim` or `sim_cap` target is used, a simulator binary, `sim`, is generated. The simulator automatically loads a physical memory image from a series of hex memory files, `mem0.hex`, `mem1.hex`, ... `mem7.hex`, used to populate initial memory contents for BRAM. By default, the memory image generated from `sw` contains a small interactive test suite that communicates via a simulated serial I/O hooked up to the simulator's standard input and output streams.


```

module ../../cherilibs/trunk/peripherals/dram.so
module ../../cherilibs/trunk/peripherals/ethercap.so
module ../../cherilibs/trunk/peripherals/uart.so

device "dram0" {
    class dram;
    addr 0x0;
    length 0x40000000;
};
#ifdef "CHERI_KERNEL" device "kernel" {
    class dram;
    addr 0x100000;
    length 0xff00000;
    option path getenv "CHERI_KERNEL";
    option type "mmap";
    option cow "yes";
};
#ifdef "CHERI_SDCARD" device "sdcard0" {
    class sdcard;
    addr 0x7f008000;
    length 0x400;
    option path getenv "CHERI_SDCARD";
    option readonly "yes";
};
#ifdef "CHERI_CONSOLE_SOCKET" device "uart0" {
    class uart;
    addr 0x7f000000;
    length 0x20;
    irq 0;
    option type "stdio";
}
#ifdef "CHERI_CONSOLE_SOCKET" device "uart0" {
    class uart;
    addr 0x7f000000;
    length 0x20;
    irq 0;
    option type "socket";
    option path getenv "CHERI_CONSOLE_SOCKET";
}

```

Figure 5.1: Example simconfig configuration file

5.6 Running the BERI test suite

The `cheritest/trunk` subtree contains a MIPS ISA unit test suite that exercises various processor features, including initial register values, memory access, jump instructions, exceptions, and so on. `make test` in the `cheri/trunk` tree will run the test suite; a detailed discussion of the test suite appears in Chapter 7.

Chapter 6

Using the BERI Debug Unit

The BERI1 prototype includes a simple debug unit that communicates with an external host using a two-way streaming 8-bit interface. The debug unit can pause and step the pipeline, set breakpoints, and insert instructions into the pipeline. Instructions inserted by the debug unit may use operands from the debug unit as well as operands from the register file. The result of an instruction from the debug unit may be written back to the debug unit as well as the register file. These debug unit elements may be used to implement a variety of higher-level services, including a proxy from the GDB server protocol, and an external memory image loader. In general, users of BERI should interact with the debug unit using `erictl` rather than interfacing directly with the `deug-unit` protocol, which is subject to change over time and between BERI versions.

6.1 Communicating with the BERI Debug Unit

The BERI debug unit communicates with a host computer over a two-way streaming 8-bit interface. The current system uses an Altera JTAG streaming component which is tunnelled over USB to the host PC. The Altera System Console utility may send and receive bytes to and from the debug unit using a desktop computer connected with a USB cable. All commands and responses are defined as series of bytes sent or received over this streaming channel.

Commands and responses traveling over the channel are arranged as messages. Every message begins with two bytes. The first is the message type and the second is the message length. If the message length is non-zero, the prescribed number of bytes follow the two bytes of the message header.

6.2 BERI Debug Registers

The BERI debug unit has eight registers:

- Debug Instruction
- Operand A
- Operand B
- Breakpoint 0

Instruction	Command	Length	Payload
Load Instruction	i	4	Instruction
Load Operand A	a	8	64-bit Operand
Load Operand B	b	8	64-bit Operand
Execute Instruction	e	0	
Report Destination	d	0	
Load Breakpoint 0	0	8	64-bit Address
Load Breakpoint 1	1	8	64-bit Address
Load Breakpoint 2	2	8	64-bit Address
Load Breakpoint 3	3	8	64-bit Address
Pause Execution	p	0	
Resume Execution	r	0	
Step Execution	s	0	
Resume Execution	r	0	
Move PC to Destination	c	0	
Resume Unpipelined	u	0	

Table 6.1: Instruction messages for the debug unit

- Breakpoint 1
- Breakpoint 2
- Breakpoint 3
- Destination

The first seven registers can be written and the Destination register can be read.

6.3 BERI Debug Instructions

The BERI debug unit supports the instructions listed in Table 6.1. The “Command” is the ASCII character that should appear in the first byte of the message sent to the debug unit, that is, the message type. The “Length” is the value of the second byte of the message. The “Payload” is the contents of the following bytes, equal in number to the value of the “Length” field. All instructions will produce a response from the debug unit confirming completion of the request. These message types are listed in Section 6.4.

Notes on Some Instructions:

Load Instruction Load an instruction into the debug unit’s Instruction register in preparation to insert it into the processor pipeline.

Load Operand A & B Load values into the Operand A & B registers for possible use as operands of the instruction in the Instruction register.

Execute Instruction Insert the instruction contained in the Instruction register into the processor pipeline.

Report Destination Report the 64-bit (8-byte) value in the Destination register. The debug unit will send a message containing the contents of the Destination register back to the debugger.

Load Breakpoint 0-3 Load an address into one of the Breakpoint registers and arm that breakpoint. When the next program counter is equal to one of the Breakpoint registers, the processor will automatically pause and the value of the breakpoint, which when fired will be sent in a message to the debugger. Loading the address `0xffffffffffffffff` will disable a breakpoint.

Step Execution Step one instruction if the processor is paused. If the next instruction is a branch, the branch delay slot and the branch target will also be executed.

Accessing Debug Registers from Debug Instructions When instructions originate from the debug unit, references to **R0** are interpreted as references to registers in the debug unit. An instruction from the debug unit which takes two operands from **R0** and writes back to **R0** will take Operand A and Operand B and will write back to the Destination register in the debug unit. In general, if “rs” refers to **R0**, that operand will come from Operand A in the debug unit and if “rt” refers to **R0**, that operand will come from Operand B in the debug unit.

6.4 BERI Debug Responses

Table 6.2 lists message types that the debug unit may generate. All of them are direct responses to instructions except for the “Breakpoint Fired” command which might be delivered at any time.

Notes for Some Responses:

Execute Instruction and Exception Responses When an “Execute Instruction” command is received, the debug unit will return an “Execute Instruction Response” message if execution of the instruction did not throw an exception. If the instruction generated an exception, the debug unit will return an “Execute Exception Response” message with a payload of one byte which will contain the 4-bit MIPS exception code generated by the instruction.

Breakpoint Fired The “Breakpoint Fired” message is sent when an instruction commits a next PC in write-back which is equal to one of the four breakpoint registers. The “Breakpoint Fired” message has a payload containing the 8-byte address value of the breakpoint that fired.

Message	Type Byte	Length	Payload
Load Instruction Response	0xe9	0	
Load Operand A Response	0xe1	0	
Load Operand B Response	0xe2	0	
Execute Instruction Response	0xe5	0	
Execute Exception Response	0xc5	1	MIPS Exception Code
Report Destination Response	0xe4	8	64-bit Value
Load Breakpoint 0 Response	0xb0	0	
Load Breakpoint 1 Response	0xb1	0	
Load Breakpoint 2 Response	0xb2	0	
Load Breakpoint 3 Response	0xb3	0	
Pause Execution Response	0xf0	0	
Resume Execution Response	0xf2	0	
Step Execution Response	0xf3	0	
Resume Execution Response	0xf2	0	
Move PC to Destination Response	0xe3	0	
Resume Unpipelined Response	0xf5	0	
Breakpoint Fired	0xff	8	64-bit Address

Table 6.2: Message types from the debug unit

Chapter 7

The BERI Unit-Test Suite

The BERI prototype includes a simple unit test suite implemented using the Python Nose framework. The test suite exercises key BERI functionality in a controlled and easily diagnosable environment, an instrumented BERI simulator, with a goal of testing both basic MIPS ISA functionality and CHERI security extensions. This chapter explains the structure and components of the test suite, how to run the test suite, and how to add new tests. It also describes some of the tools available for diagnosing test results.

7.1 The BERI Unit-Test Environment

The BERI unit test suite is implemented using a combination of the BERI Bluespec simulator (with extensions for debugging), `make`, the MIPS toolchain, the Python Nose test framework, and a moderate collection of test programs and Nose classes to evaluate test output. The unit test suite can also be run against the `gxemul` MIPS simulator, which has proven useful for checking our interpretation of the MIPS ISA against a more common interpretation. In the future, we hope also to run the test suite against BERI synthesised in an FPGA, likely with the help of JTAG.

7.2 Software Dependencies

To run the BERI unit test suite, you will need to have the following software installed:

- Python
- The MIPS SDE cross-linker, installed as *sde-ld*.
- The CTSRD-modified version of the GNU binutils, available from <https://github.com/CTSRD-CHERI/binutils>. The assembler should be installed as *mips64-as*, objcopy as *mips64-objcopy* and so on.

If you want to run the tests that are written in the C language with capability extensions, you will also need the CTSRD-modified version of the clang compiler, available from <https://github.com/CTSRD-CHERI/llvm>.

If you want to run the tests against GXEMUL, you will also need the CTSRD-modified version of GXEMUL, available from <https://github.com/CTSRD-CHERI/gxemul>. The modified GXEMUL does not

Directory	Description
<code>cheritest/trunk/</code>	Root of the BERI test suite tree, home of the makefile, linker scripts, and test library code
<code>gxemul.log</code>	Destination for gxemul test run output
<code>log/</code>	Destination for BERI simulator test run output
<code>obj/</code>	Destination for test object files, memory images, and assembly dumps
<code>tests/</code>	Various subdirectories holding source code for individual tests, and their matching Python Nose classes
<code>tools/</code>	Utility functions to perform common functions such as interpreting BERI simulator and gxemul output

Table 7.1: Directories in the BERI unit test suite

include support for the capability instructions, but it does include modifications to integrate it with our test framework, and improved emulation of floating point instructions.

If you want to run the tests against the formal model of MIPS developed by the REMS (“Rigorous Engineering for Mainstream Systems”) project, you will also need their MIPS formal model, available from <http://www.cl.cam.ac.uk/acjf3/I3>. You will also need REMS project’s support libraries to turn the model into an executable specification; at the time of writing, these had not been released.

If you want to run the tests against a Bluespec-level simulation of BERI, you will also need the Bluespec tools, and a compatible version of GCC. (Bluespec is compiled into C++ which is then compiled by g++).

7.2.1 BERI Test-Suite Directory Layout

Table 7.2.1 describes the directories in the BERI unit test suite.

7.2.2 BERI ISA Extensions for Testing

The BERI test suite employs debugging extensions to the 64-bit MIPS ISA to examine the state of a simulated BERI system after each test. It dumps the general-purpose register file, the CP0 registers and the capability coprocessor registers, and allows tests to terminate the simulation in a controlled manner. Current extensions are exposed via CP0 register operations, as shown in table 7.2. In the future they will likely move to capability coprocessor extensions to reduce the possibility of a collision with the existing MIPS ISA. In the future, we anticipate the addition of further extensions in support of testing to dump the simulation memory image.

7.2.3 Unit Test Support Library

Most BERI unit tests are linked against a thin loader, `init.s`, which is responsible for setting up various aspects of CPU and memory configuration. They serve to

Instruction	Description
<i>mtc0 register</i> , \$26	Dump arithmetic registers to trace
<i>mtc2 register</i> , \$0, 6	Dump capability registers to trace
<i>mtc0 register</i> , \$23	Stop the simulation

Table 7.2: BERI ISA extensions for testing

- Set up a stack at the top of memory.
- Install default before- and after-boot exception vectors and handlers, which will dump the register file and terminate if triggered.
- Explicitly clear all general-purpose registers except stack-related registers that may have been modified during startup.
- Invoke a user-provided `test` function using `JAL`; currently all test functions are implemented in assembly, but the calling convention should support C as well.
- On return from `test`, dump the register file and terminate.

In addition, a small library of support routines common to more complex tests, including functions for copying memory and installing exception handlers, may be found in `lib.s`. We anticipate that this library will grow in size as the test suite is made more comprehensive.

A few low-level tests, referred to as *raw tests*, execute directly rather than via `init.s`, and are not linked against `lib.s`. Raw tests perform low-level verification of CPU functionality required to reliably run `init.s`, such as initial register file values on CPU reset, arithmetic instructions, the reliability of branch and jump instructions, and basic memory operations. Whenever possible, writing raw tests should be avoided, because they necessarily replicate functionality (such as register dumping), and lack access to a pre-configured stack.

Note that all tests will be run twice by the suite – once from uncached instruction memory, and once from cached instruction memory. Timing and pipeline effects differ significantly between the two cases. One impact of this difference is that all tests must be relocatable and able to run in multiple MIPS `xkphys` segments.

7.3 Running the BERI Test Suite

Typically, the test suite will be run as follows:

```
$ cd ctsrd/cheritest/trunk
$ make test
```

The `CHERIROOT` variable may be used to tell the test suite where to find BERI tools for processing memory images and the BERI simulator; the BERI simulator must first have been built using `make sim` or similar. The test suite may be run against `gxemul` as follows:

```
$ cd ctsrd/cheritest/trunk
$ make gxemul-build
$ make gxemul-nosetest
$ make gxemul-nosetest_cached
```

7.3.1 Jenkins

If you are developing in the Cambridge development environment, the BERI unit test suite is run automatically by the Jenkins build framework. Jenkins can be monitored by visiting the following URL:

<https://ctsrd-build.cl.cam.ac.uk>

7.4 Unit-Test Structure

Each unit test consists of a short assembly program that exercises specific features in the BERI CPU, and a Nose class that contains a set of assertions about termination state for the test. Modifications to the test suite typically take the form of modifying an existing test to check new assertions, or adding an entirely new test via a new test program and set of corresponding assertions.

7.4.1 Test Types

Tests are split into two categories: raw tests that have few low-level dependencies and are intended to exercise basic CPU features such as the register file, and higher-level tests that are able to depend on common CPU initialization code and a support library. Raw tests are necessarily run before higher-level tests, which typically depend on features checked in raw tests. Raw test files are prefixed with `raw_`, and higher-level test file names are prefixed with `test_`; the build framework uses these prefixes to identify assembly and linking requirements, so they must be used.

Unless there is a specific reason to do so, new tests should be added as higher-level tests, relying on the `init.s` framework to set up the stack, dump register state on completion, and terminate the simulator, rather than hand-crafting this code. This provides access to routines such as `memcpy` that are frequently useful when implementing tests.

7.4.2 Test Structure

All tests are compiled using 64-bit MIPS instructions, and attempt to follow a standard application binary interface (ABI) to allow easy reuse of compiled MIPS code reused in the test environment. Currently, no C code is linked into the test suite, but it is easy to imagine doing so in the future, making ABI conformance critical.

High-level tests implement a single, global function `test`. When `test` terminates, the calling code in `init.s` will dump register state and terminate the simulator; these registers then become available to the Nose test framework for checking. Other than changes to the program counter, `$PC`, the test framework avoids any changes to register values after the test returns. Tests may rely on the availability of a roughly 1K stack. Tests execute in both the cached but unmapped region of memory around `0x9800000040000000`, with a stack growing down from `0x9800000080008000`, and the uncached and unmapped region around

0x9000000040000000, with a stack growing down from 0x9000000080000000, but may make use of any required processor features such as cached and mapped memory regions, CP0 MMU operations, etc.

7.4.3 Test Termination

Normally, high-level tests will terminate by returning from the `test` function, triggering a register dump and simulator termination. However, the test framework is executed with a 100000-cycle limit on simulation time in order to ensure termination, catching (for example) infinite loops in software, or exception cycles. As tests become more complicated, this limit may need to be changed, but is present to ensure that tests will eventually always terminate, even if software enters an infinite loop.

7.4.4 Connecting New Tests to the Build

Nose test files must begin with the prefix `test_`, which will normally occur for high-level tests; Nose test files for raw tests will therefore be prefixed with `test_raw_`. New unit tests are hooked up to the build system by adding their source files to the `TEST_FILES` variable in the makefile. This is normally done by adding the test filename to one of the make variables for a test subset such as `TEST_ALU_FILES`. For the time being, all test source and Nose files must be placed under the `tests` directory in an appropriate sub-directory which should be included in the `TESTDIRS` variable.

7.4.5 Test Attributes

Each of the Python test scripts is tagged with Python attributes that indicate which versions of BERI (or generic MIPS) the test is expected to run on. The default (no attributes) is that the test is expected to work on any processor that complies with the MIPS R4000 specification. Dependences on additional features are indicated as follows:

7.5 Example Unit Test: Register Zero

To explore the above design, we will consider the `test_reg_zero` unit test, which checks that the MIPS general-purpose register **R0**, also known as **\$zero**, has the required special property that it always return the value 0. The correct functioning of **\$zero** is not required for any raw tests, nor `init.s`, so the test is placed in the high-level test suite. The test performs a number of activities:

- Sets up a stack from for the function `test` by manipulating **\$sp** and **\$fp**.
- Pushes the return address, **\$ra**, and saved frame pointer, **\$fp**, onto the stack.
- Copies a value from **\$zero** into **\$t0** for inspection.
- Assigns a value to **\$zero** from an immediate, and then copies out to **\$t1** to confirm that the value does not get saved.
- Assigns a value to **\$zero** from a register, and then copies out to **\$t2** to confirm that the value does not get saved.
- Restores **\$fp** and **\$ra** from the stack and returns.

Attribute	Description
beri	Test depends on BERI implementation details
beriinitial	Initial values of registers same as BERI
cached	Test must be run from cached memory
capabilities	CHERI capability unit
comparereg	CP0 Compare register
config3	CP0 Config3 register
counterdev	Counter device
dumpicache	
float	Floating point unit
float32	Floating point unit that supports 32-bit mode
float64	Floating point unit that supports 64-bit mode
floatcmov	Floating point conditional move instructions
floatexception	Floating point unit can raise exceptions
floatflags	FPU supports IEEE condition flags
floatfcsr	
floatfexr	
floatfenr	
floatindexed	
floatpaired	Floating point unit that supports paired single
floatrecip	RECIP.D and RECIP.D instructions
floatrsqrt	RSQRT.S and RSQRT.D instructions
ignorebadex	32-bit arithmetic ignores the top 32 bits
invalidateL2	cache instruction for L2 cache
llsc	Load-linked and store conditional instructions
llscnotmatching	SC will fail if address doesn't match LL
llscspan	SC succeeds even if there is a load after LL SC fails if there is a store after LL
lladdr	CP0 LLAddr register
rdhwr	rdhwr instruction
swi	Software interrupts
tlb	Translation lookaside buffer
smalltlb	A TLB just like CHERI2's
bigtlb	A TLB just like CHERI1's
gxemultlb	A TLB just like GXEMUL's
largepage	
mt	Multi-threaded CPU
mtc0signex	mtc0 sign-extends the value moved
nofloat	Test will only work if FPU is absent
trapi	TRAPI instruction
userlocal	User local register
watch	Watch points

Table 7.3: Test attributes

7.5.1 Register Zero Test Code

Example assembly source code is illustrated in Figure 7.5.1.

7.5.2 Register Zero Nose Assertions

Figure 7.5.2 illustrates the Nose assertion set for this test, confirming a number of desired properties that should hold after the test code runs:

- that **\$zero** held zero on exit,
- that **\$t0** held zero on exit, meaning that a simple move from **\$zero** held zero on start,
- and that registers **\$t1** and **\$t2** held zero values, meaning that various writes to **\$zero** did not change the value returned when reading the register.

7.6 Conclusion

This chapter has introduced the BERI unit test suite, exploring both the structure of the suite and the implementation of individual tests. The test suite is intended to supplement formal methods by testing the programmer-level view of ISA correctness, and while it cannot be authoritative regarding the correctness of BERI, it is extremely valuable in development, exercising critical instruction combinations and providing clear diagnostics. We hope to introduce a new unit test for each bug encountered in BERI, and expand the test suite to provide detailed coverage of new ISA features.

```

.set mips64
.set noreorder
.set nobopt
.set noat

#
# This test checks that register zero behaves the way it should: each of
# $t0, $t1, and $t2 should be zero as at the end, as well as $zero.
#

.global test
test: .ent test
daddu $sp, $sp, -32
sd $ra, 24($sp)
sd $fp, 16($sp)
daddu $fp, $sp, 32

# Pull an initial value out
move $t0, $zero

# Try storing a value into it from an immediate
li $zero, 1
move $t1, $zero

# Try storing a value into it from a temporary register
li $t3, 1
move $zero, $t3
move $t2, $zero

ld $fp, 16($sp)
ld $ra, 24($sp)
daddu $sp, $sp, 32
jr $ra
nop # branch-delay slot
.end test

```

Figure 7.1: Example regression test checking properties of **\$zero**

```

from beritest_tools import BaseBERITestCase

class test_reg_zero(BaseBERITestCase):
    def test_zero(self):
        '''Test that register zero is zero'''
        self.assertRegisterEqual(self.MIPS.zero, 0,
            "Register zero has non-zero value on termination")

    def test_t0(self):
        '''Test that move from zero is zero'''
        self.assertRegisterEqual(self.MIPS.t0, 0,
            "Move from register zero non-zero")

    def test_t1(self):
        '''Test that immediate store of non-zero to zero returns zero'''
        self.assertRegisterEqual(self.MIPS.t1, 0,
            "Immediate store to register zero succeeded")

    def test_t2(self):
        '''Test that register store of nonzero to zero returns zero'''
        self.assertRegisterEqual(self.MIPS.t2, 0,
            "Register move to register zero succeeded")

```

Figure 7.2: Example Nose assertion file for the **\$zero** test

Chapter 8

BERI on Altera FPGAs

WARNING: the contents of this chapter are frequently out of date as a result of evolving Altera tools and changes to the BERI build environment.

This chapter describes how to build BERI for synthesis using Bluespec, configure the Altera build environment, and synthesize BERI for the Terasic tPad and DE4 FPGA teaching boards described in later chapters. This information is relevant to researchers working with the BERI hardware design. Software consumers of BERI can find information on using specific Terasic boards in Chapter 9, and will not need to follow the directions in this chapter.

8.1 Building BERI for Synthesis

BERI source code may be compiled to Verilog with the verilog target.

```
$ CAP=1 make verilog
$ ./sim
```

The BERI Verilog build is also sensitive to five `make` variables described in Section 5.3. The result of building BERI for synthesis is a set of Verilog files in the appropriate directory in `ip/`, with the file `mkTopAvalonPhy.v` containing the top-level module. These files may be copied into one of the directories in the `boards/` directory to be synthesized for a particular board. As a note, when a supported FPGA board is connected to the computer by USB, it is possible to connect to the BERI UART using JTAG:

```
$ nios2-terminal --instance 0
```

8.2 The Altera Development Environment

Terasic's FPGA evaluation boards include Altera FPGAs; the following sections depend on the correct installation of Altera's FPGA development toolchain in order to synthesize and program the on-board FPGAs. Some of Altera's tools – especially the GUIs, but also some command-line tools – require X11; in these cases, if using a central build server, ensure that the `-X` argument is passed to the `ssh` command:

```
$ ssh -X user@zenith.cl.cam.ac.uk
```

To configure your shell to use Bluespec, Altera, and other development toolchain elements for BERI, such as compilers and linkers, use the following script from the CTSRD Subversion repository (described in previous chapters):

```
$ cd ctsrd/cheri/trunk
$ source setup.sh
```

In order to successfully build the BERI hardware project, ensure that you have added the Bluespec Verilog library to the Quartus global library path. This library is typically located at:

```
<BluespecDirectory>/lib/Verilog
```

Also ensure that you have added any relevant license files needed to build the project. For example, if you are using an Terasic touchscreen, you may need to add the license file for the `i2c_touch` Verilog module to the license file string for Quartus.

The Quartus 12.1 tools assume that `/bin/sh` is a link to `/bin/bash`, but on recent Ubuntu systems the symbolic link is to `/bin/dash`. This inconsistency breaks some of the Qsys generate scripts. There are two solutions:

1. Patch the Quartus 12.1 distribution so that scripts starting:

```
#!/bin/sh
```

```
now start
```

```
#!/bin/bash
```

(The Computer Lab uses this approach, and recommends it.)

2. Modify the symbolic link so that `/bin/sh` links to `/bin/bash`. Although we've never found a problem with this solutions, it seems inelegant and runs the risk of breaking some aspect of your Linux setup.

Finally, if you are using Ubuntu, you may need to insert a new rules file into `/etc/udev/rules.d/`. You might add a new file named `51-usbblaster.rules` with the following contents:

```
# Set permissions for Altera USB Blaster
SUBSYSTEM=="usb", ATTR{idVendor}=="09fb", ATTR{idProduct}=="6001", \
MODE="0666", OWNER="root", GROUP="dialout"
# Set permissions for Fast Altera USB2 Blaster
SUBSYSTEM=="usb", ATTR{idVendor}=="09fb", ATTR{idProduct}=="6810", \
MODE="0666", OWNER="root", GROUP="dialout"
```

Directory	Board
<code>ctsrd/cheri/trunk/boards/terasic_tPad</code>	Terasic tPad
<code>ctsrd/cheri/trunk/boards/terasic_de4</code>	Terasic DE4

Table 8.1: Terasic per-board directories

Target	Description
<code>all</code>	builds everything using the following steps (except download)
<code>build_cheri</code>	builds the BERI processor
<code>build_peripherals</code>	builds the peripherals
<code>build_miniboot</code>	builds miniboot ROM and copy initial.hex here
<code>build_qsys</code>	builds Qsys project containing BERI, etc.
<code>build_fpga</code>	synthesize, map, fit, analyze timing, and generate FPGA image
<code>report_critical</code>	scans <code>build_fpga</code> reports for critical warnings
<code>report_error</code>	scans <code>build_fpga</code> reports for errors
<code>download</code>	attempts to download the FPGA (.sof) image to the FPGA but the chain file (.cdf) may need to be updated for your configuration (e.g. USB port number)
<code>clean</code>	removes Quartus and Qsys build files
<code>cleanall</code>	clean + clean peripherals, BERI and miniboot

Table 8.2: Make targets for per-board directories

8.3 Synthesizing BERI

WARNING: This section applies to the DE4 board setup. The tPad setup has atrophied and needs further attention.

The CTSRD project provides reference configurations for BERI on the Terasic tPad and DE4 boards; per-board directories are listed in Table 8.1. Each board directory contains its own Quartus project, `Makefile`, etc. Table 8.2 shows the available make targets.

Targets `build_cheri` and `build_peripherals` cause other Makefiles to be used to build various Verilog components that are found by Quartus via the paths in `peripherals.ipx` and `processors.ipx`. `build_miniboot` compiles the miniboot loader C code and produces a ROM image `initial.hex`, which is copied into the board directory.

A `make cleanall;make all` will take around 40 minutes to an hour to complete. Then, to download the system to an FPGA board plugged into the host PC, `make download` may work – although a new chain (.cdf) file might be needed to reflect your particular setup. This result is most easily achieved using the GUI in Quartus to open the programmer and configure.

Chapter 9

BERI on Terasic tPad and DE4

This chapter describes how to use the BERI processor prototype on the Terasic tPad and DE4 FPGA teaching boards. It includes tutorial material on programming the board and on how board peripherals are exposed to BERI in the reference designs provided by the CTSRD project. This chapter is intended to support software development on BERI. See Chapter 8 for documentation pertinent to hardware development. Later chapters describe building and using CheriBSD on Terasic boards.

9.1 BERI Configuration on the tPad and DE4

Communication with external I/O devices, such as NICs, is accomplished via a blend of memory-mapped I/O, interrupts, and (eventually) DMA. The BERI processor and operating system stack supports a variety of peripherals ranging from Altera “soft” cores, such as the JTAG UART and SD Card IP cores, to “hard” peripherals provided by Terasic on its tPad and DE4 development boards. The following sections document available peripherals and their configuration on the Avalon system-on-chip bus as configured in the BERI reference designs.

9.1.1 Physical Address Space on the tPad

Table 9.1 shows the physical addresses reserved for I/O devices in the BERI reference tPad configuration.

9.1.2 Physical Address Space on the DE4

Table 9.2 shows the physical addresses reserved for I/O devices in the BERI reference DE4 configuration.

9.2 Altera IP Cores

BERI and FreeBSD support a number of Altera “soft” IP cores on the Terasic tPad and DE4 platforms. Many of these IP cores are documented in the *Embedded Peripherals IP User Guide*¹ provided by Altera, including the JTAG UART core and Avalon-MM and Avalon-ST bus attachments.

¹http://www.altera.com/literature/ug/ug_embedded_ip.pdf

Base address	Length	IRQ	Description
0x7f000000	64	0	Altera JTAG UART
0x7f001000	64	-	Altera JTAG UART for debugging output
0x7f002000	64	-	Altera JTAG UART for data
0x7f004000	4	-	Old location of count register until 2013-03-01
0x7f007000	1024	-	Altera Triple-Speed Ethernet MegaCore MAC control
0x7f007400	8	-	... MAC transmit FIFO
0x7f007420	32	2	... MAC transmit FIFO control ¹
0x7f007500	8	-	... MAC receive FIFO
0x7f007520	32	1	... MAC receive FIFO control ²
0x7f008000	1024	3	Altera University Program Secure Data Card IP Core
0x7f00A000	20	-	Hardware Version ROM ³
0x7f800000	8 MB	-	Bluespec Peripheral Address Space
0x7f800000	8	-	... Count Register (from 2013-03-01)
0x7f804000	16 KB	-	... BERI PIC_CONFIG_BASE
0x7f804000	16 KB	-	... BERI PIC_CONFIG_BASE_0, In a dual core system - Core 0 PIC
0x7f808000	16 KB	-	... BERI PIC_CONFIG_BASE_1, In a dual core system - Core 1 PIC

¹ See “Avalon-MM Write Slave to Avalon-ST Source”

² See “Avalon-ST Sink to Avalon-MM Read Slave”

³ See Table 9.3

Table 9.1: Bus configuration for BERI’s reference tPad configuration

Base address	Length	IRQ	Description
0x70000000	128 MB	-	Cambridge Multitouch LCD + 256 Mb Intel StrataFlash
0x7f000000	64	0	Altera JTAG UART
0x7f001000	64	7	Altera JTAG UART for debugging output
0x7f002000	64	8	Altera JTAG UART for data
0x7f002100	32	6	Bluespec 16550 UART (RS232) ⁰
0x7f004000	4	-	Old location of count register until 2013-03-01
0x7f005000	1024	-	Altera Triple-Speed Ethernet MegaCore MAC control Port 1
0x7f005400	8	-	... MAC transmit FIFO
0x7f005420	32	11	... MAC transmit FIFO control ¹
0x7f005500	8	-	... MAC receive FIFO
0x7f005520	32	12	... MAC receive FIFO control ²
0x7f006000	1	-	DE4 LEDs, one bit per LED
0x7f007000	1024	-	Altera Triple-Speed Ethernet MegaCore MAC control Port 0
0x7f007400	8	-	... MAC transmit FIFO
0x7f007420	32	2	... MAC transmit FIFO control ¹
0x7f007500	8	-	... MAC receive FIFO
0x7f007520	32	1	... MAC receive FIFO control ²
0x7f008000	1024	-	Altera University Program Secure Data Card IP Core
0x7f009000	2	-	Switches and Buttons one bit each (DIP[0:7], SW[0:3], BUTTON[0:3])
0x7f00A000	20	-	Hardware Version ROM ³
0x7f00B000	8	9	OpenCores i2c Controller for the HDMI chip
0x7f00B080	1	-	1-bit PIO to reset the HDMI chip
0x7f00C000	8	-	Temperature and Fan Control ⁴
0x7f100000		-	Philips ISP1761 USB 2.0 Chip ⁵
0x7f100000	256 KB	5	... Host Controller
0x7f140000	4	4	... Peripheral Controller
0x7f800000	8 MB	-	Bluespec Peripheral Address Space
0x7f800000	8	-	... Count Register (from 2013-03-01)
0x7f804000	16 KB	-	... BERI PIC_CONFIG_BASE
0x7f804000	16 KB	-	... BERI PIC_CONFIG_BASE_0, In a dual core system - Core 0 PIC
0x7f808000	16 KB	-	... BERI PIC_CONFIG_BASE_1, In a dual core system - Core 1 PIC

⁰ See Section 9.6

¹ See “Avalon-MM Write Slave to Avalon-ST Source”

² See “Avalon-ST Sink to Avalon-MM Read Slave”

³ See Table 9.3

⁴ See Section 9.5

⁵ See Philips ISP1761 Hi-Speed Universal Serial Bus On-The-Go controller

Base	Length	Item	Format
0x0	4 Bytes	Build Date	Binary coded decimal formatted as mmddyyyy
0x4	4 Bytes	Build Time	Binary coded decimal formatted as 00hhmmss
0x8	4 Bytes	Svn Version	Binary coded decimal
0xc	8 Bytes	Host Name	ASCII, truncated to 8 characters

Table 9.3: Contents of the BERI Hardware Build Version Number ROM

Certain Altera IP cores are described in other documents, including the Altera Triple-Speed MAC described in the *Triple-Speed Ethernet MegaCore Function User Guide*², and SD Card IP core described in the *Altera University Program Secure Data Card IP Core*³ documents from Altera.

9.3 Cambridge IP Cores

Cambridge provides two “soft” peripheral devices: the *count device*, which simply provides a memory-mapped register that is incremented on every read (intended for cache testing), and a memory-mapped interface to the Terasic MTL multitouch LCD panel. This latter IP core includes both memory-mapped support for a pixel frame and a VGA-like text frame buffer suitable for use as a system console. It also provides access to multitouch input.

9.3.1 The DE4 Multitouch LCD

Hardware Overview

A Terasic MTL-LCD is connected to the DE4 via the supplied ribbon cable. This connection provides a parallel interface running at 33 MHz to drive the LCD and an I2C interface to obtain touch information. Terasic provides an encrypted block (`i2c_touch_config.v`) to talk I2C to the touch panel and exports parameters as a simple parallel interface.

We have built three key hardware components to interface to the MTL-LCD:

MTL_LCD_Driver – This peripheral takes an AvalonStream of pixel values and maps them to the MTL (multi-touch) LCD color screen, which has an 800x480 resolution. Pixels are 24-bits (8-bit red, green, blue). The main clock must run at the pixel clock rate of 33 MHz. The clock to the MTL-LCD (`mtl_dclk`) must be fed to the LCD outside of this module. A dual-clock FIFO is needed in the AvalonStream between this module and the `MLT_Framebuffer_Flash`.

²http://www.altera.com/literature/ug/ug_ethernet.pdf

³ftp://ftp.altera.com/up/pub/Altera_Material/11.0/University_Program_IP_Cores/Memory/SD_Card_Interface_for_SoPC_Builder.pdf

MTL_LCD_HDMI – This peripheral is an alternative to MTL_LCD_Driver which runs the multitouch LCD out of spec. (but still working just fine) in order to mirror to HDMI (and via HDMI to VGA) at 720x480 pixels with the correct timing specification. H-sync and V-sync timings are changed and the pixel clock is reduced to 27MHz. This reduced pixel clock rate has the advantage that the bandwidth from the SSRAM frame buffer memory is less demanding. As with the MTL_LCD_Driver, this module is connected via a dual-clock FIFO and an AvalonStream interface to the MTL_Framebuffer_Flash. No changes are needed to MTL_Framebuffer_Flash to use this module since the difference in pixel clock rate is accommodated by flow-control in the AvalonStream.

MTL_Framebuffer_Flash – This component provides a memory-mapped frame buffer using the DE4's off-chip SSRAM to store the frame buffer and provides access to the Flash, which is on the same bus as the SSRAM. It provides an Avalon memory-mapped interface that allows a processor to write to the SSRAM. This module is designed to work at the main system clock rate of 100 MHz. Note that the clock to the SSRAM needs to be provided outside of this module, direct from a PLL. The SSRAM conduit interface must be connected to the SSRAM pins. The I2C conduit interface (coe_touch) must be connected to Terasic's I2C encrypted block outside of the Qsys project.

In addition, the following libraries of ours are used:

AlteraROM provides a font ROM initialized from `fontrom.mif`.

VerilogAlteraROM.v provides Verilog wrapped by AlteraROM.

Avalon2ClientServer provides the Avalon memory-mapped. interface

AvalonStreaming provides the Avalon streaming interface.

Software Overview

The MTL_Framebuffer is accessed via an 8 MB memory mapped region where the first 2 MB maps the SSRAM, which contains both the text and pixel frame buffer. Control registers start 4 MB into the region. Random access reads and writes of arbitrary size are permitted to the main frame buffer, but registers require 32-bit accesses. Note that writes to the frame buffer are queued and incur little latency whereas reads need to schedule access around the LCD updates so incur a much greater latency penalty. Reads and writes to registers are quick.

The pixel frame buffer is 32 bits per pixel. The upper byte is ignored, but followed by bytes of red, green and blue channels. The resolution is 800x480 with the first pixel being top level. The text frame buffer accepts characters of 16-bits with the upper byte representing the VGA text color and the lower byte holding the character. There are 100 columns and 40 rows of text. VGA text color is a byte in the following format: (1-bit flashing, 3-bit background color, 4-bit foreground color). Colors are from the following table:

code	color	code	color
0	black	8	dark grey
1	dark blue	9	light blue
2	dark green	10	light green
3	dark cyan	11	light cyan
4	dark red	12	light red
5	dark magenta	13	light magenta
6	brown	14	light yellow
7	light grey	15	white

See `mtl_test_small.c` for an example which drives the MTL-LCD using a NIOS for some helper functions, and so on. Table 9.4 describes the memory map of the MTL-LCD.

The frame-buffer-blending register has the following format (from MSB to LSB):

- Top 4 bits are unused, but should be set to zero.

base offset	length	description
0x0000000	2MB	SSRAM
0x0000000	800x480 words	pixel frame buffer, 32-bit color, although only 24 bits used: 8 bits each of (r,g,b) where b is the LSB
0x0177000	100x40x2 bytes	of text buffer in its default location
0x0400000	1 word	frame buffer blending (see below)
0x0400004	1 word	text cursor position, bytes: (unused, unused, x, y)
0x0400008	1 word	character frame buffer offset base address relative to the start of the SSRAM, 0x177000 after reset (i.e., 800x480 words, so just after the pixel buffer). Note that this must be a 32-bit word aligned offset.
0x040000c	1 word	touch point x1, -1 if not valid
0x0400010	1 word	touch point y1, -1 if not valid
0x0400014	1 word	touch point x2, -1 if not valid
0x0400018	1 word	touch point y2, -1 if not valid
0x040001c	1 word	(touch_count,gesture), -1 if not valid, where touch_count is a 2-bit value (0,1 or 2 touches) and gesture is an 8-bit value (see table below for details). Reading this register dequeues all of the current touch sensor values.
0x4000000	64MB	Flash memory (see below)

Table 9.4: Memory map used for the MTL device

- 4 bits of VGA color code providing a default color for the whole screen. After reset, this is set to 2 (dark green). Typically, this will need to be set to 0 for general use.
- 8 bits of alpha blending for the pixel frame buffer. This value is subtracted using saturation arithmetic from the character colors; a value of 255 erases the character frame buffer. Reset value is 255 (characters off).
- 8 bits of alpha blending for the character frame buffer foreground color, subtracted from the pixel background color using saturation arithmetic. 255 makes the characters opaque, and 0 makes them transparent. Reset value is 255 (pixel off) when the character pixel is on.
- 8 bits of alpha blending for the character frame buffer background color, which is subtracted from the pixel color using saturation arithmetic. 255 makes the background opaque, and 0 makes the background transparent. Set to 255 after reset (pixel off) when the character pixel is off.

The MTL two-touch gesture codes (copied from the MTL-LCD manual):

code	gesture
0x30	north
0x32	north-east
0x34	east
0x36	south-east
0x38	south
0x3A	south-west
0x3C	west
0x3E	north-west
0x40	click
0x48	zoom in
0x49	zoom out

9.3.2 HDMI Chip Configuration via I2C

We use the Terasic HDMI_TX_HSMC daughter card on the DE4 board to obtain HDMI output mirroring. Pixel data, H-sync and V-sync are provided by MTL_LCD_HDMI (see Section 9.3.1) when mirroring the multitouch LCD. However, to obtain output, the HDMI chip on the daughter card must be configured via an I2C interface. To do this, we use the I2C master interface from OpenCores⁴. This interface is wrapped in an Avalon interface that we have written⁵, which is colocated with documentation⁶

Currently the miniboot uses this I2C interface to configure the HDMI chip to a fixed output resolution. The code was hacked up in time for the November 2012 PI meeting, and will need further work when we want to dynamically change the display resolution and other details. But first we'll need a frame buffer that can produce programmable resolutions.

⁴<http://opencores.org/project,i2c>

⁵[cherilibs/trunk/peripherals/i2c/i2c_avalon.sv](#)

⁶[cherilibs/trunk/peripherals/i2c/i2c_rev03.pdf](#)

9.4 Standalone HDMI Output

The standalone HDMI output (HDMI_Driver) is an alternative to the mirrored HDMI output from the MTL-LCD discussed in Section 9.3.2. The motivation is to provide support for video streams of different resolutions from other sources (e.g. streaming out of high-bandwidth memory like DDR2 memory).

In order to support multiple resolutions, a variable pixel clock is required (Section 9.4.1) together with a software configurable HDMI timing generator (Section 9.4.2) and the HDMI chip configuration via I2C discussed earlier in Section 9.3.2. Note that we currently use the I2C interface to place the HDMI chip into DVI compatibility mode. In this mode, the resolution can be set by changing the pixel clock frequency and video timing (sync signals) without further configuration of the HDMI chip. The HDMI chip documentation is so poor that it's difficult to determine whether this is the correct usage, but it appears to work.

9.4.1 Reconfigurable Video Pixel Clock

This is a simple Qsys peripheral written in SystemVerilog to provide an Avalon memory mapped interface to Altera provided reconfigurable PLL. The reconfigurable PLL needs to be instantiated outside of this module using an ALTPLL megafunction with its reconfiguration interface enabled.

Inside this peripheral, an ALTPLL_RECONFIG is instantiated that provides a cache of the PLL parameters and, when triggered, writes them to the ALTPLL using a proprietary serial interface. ALTPLL_RECONFIG also resets the ALTPLL post configuration.

This module is addressed as follows. **All addresses refer to 32-bit little-endian words. Byte addressing is not supported.**

The lower address bits have the following meaning:

- bits 1-0 are always zero (word aligned)
- bits 5-2 is the counter_type
- bits 8-6 is the counter_parameter
- bit 9 When=0 it refers to the ALTPLL_RECONFIG parameters (above). When =1 for a write it causes the PLL parameters to be written to the PLL. When =1 and reading, it returns busy=-1, done=0.

counter_type and counter_parameters are defined in Altera's ALTPLL_RECONFIG Users Guide⁷ with the parameters for Stratix IV PLLs appearing on pages 45–46.

For Stratix IV parts (e.g. on the DE4 board) the following counter_parameters are particularly useful:

counter_parameter number	variable name	meaning
0	n	master multiplier
1	m	master divisor
4	c0	further divisor for clock 0

The output frequency clock c0 is given by: $f_{out_c0} = (n \times f_{in}) / (m \times c0)$

Where f_{out_c0} is the output frequency for clock 0 on the PLL and f_{in} is the input clock frequency (typically from an external pin on the DE4 board running at 50MHz).

For each of these counter_parameters, the following counter_types need to be set (e.g. for a required value v where $v > 0$):

⁷http://www.altera.co.uk/literature/ug/ug_altpll_reconfig.pdf

counter_type number	variable name	bit width	value from v
0	high_count	(9-bits)	$(v+1)/2$
1	low_count	(9-bits)	$v - \text{high_count}$
4	bypass	(1-bit)	$(v==1) ? 1 : 0$
5	odd_count	(1-bit)	$v \& 0x1$

After setting the above parameters in the ALTPLL_RECONFIG cache, you need to trigger the cache to write the parameters to the ALTPLL by writing some word of data (the data is irrelevant) to an address on this peripheral with address bit 9 set.

9.4.2 HDMI Timing Driver

The Qsys peripheral (HDMI_Driver) takes an AvalonStream of pixel values and maps them to the Terasic HDMI Transmitter daughter card (HDMI_TX_HSMC). It needs to be clocked at the video pixel clock frequency which may be variable. Thus, an Avalon clock crossing bridge is needed to interface to the AvalonMM slave interface which allows the following parameters to be set from software.

Address map (32-bit word offset, little-endian 12-bit values in 32-bit word)		
0	x-resolution	(in pixels)
1	horizontal pulse width	(in pixel clock ticks)
2	horizontal back porch	(in pixel clock ticks)
3	horizontal front porch	(in pixel clock ticks)
4	y-resolution	(in pixels/lines)
5	vertical pulse width	(in lines)
6	vertical back porch	(in lines)
7	vertical front porch	(in lines)

9.5 Temperature and fan control

The temperature and fan control peripheral has two read-only 32-bit registers. The first (address 0x0) returns the last temperature reading as a 32-bit signed integer in degrees Centigrade. The second (address 0x4) is the power to the fan as a range from 0 to 255.

9.6 Bluespec 16550 UART

We have created a portable Bluespec RS232 UART, based on the widely used NS16550 interface, for use with BERI. The UART needs to be accessed using byte accesses. Each of the byte-wide registers must be on 32-bit boundaries. This device is little endian. The table below lists the register with 32-bit word address offsets.

Address map (32-bit word offset, little-endian 8-bit values in 32-bit word)		
Offset	Register	Description (R=read, W=write)
0	UART_DATA	TX (W) and RX (R) buffers
1	UART_INT_ENABLE	interrupt enable (RW)
2	UART_INT_ID	interrupt identification (R)
2	UART_FIFO_CTRL	FIFO control (W)
3	UART_LINE_CTRL	Line Control Register (RW)
4	UART_MODEM_CTRL	Modem Control (W)
5	UART_LINE_STATUS	Line Status (R)
6	UART_MODEM_STATUS	Modem Status (R)
7	UART_SCRATCH	scratch register (RW)

Note: the UART_SCRATCH register can be used, for example, to check that read and write access is possible to this peripheral.

9.7 Terasic hard peripherals

9.7.1 Intel StrataFlash 64M NOR flash

The DE4 board has a single Intel StrataFlash embedded memory. Cambridge has the part with 64 MB (512 Mb), which is 16 bits wide. Note that this part might be in one package, but it has two die-stacked internal flash chips that work independently. This flash memory sits on the same bus as the SSRAM used for the frame buffer; the memory transactions are handled by the multitouch display hardware.

Read mode

After reset, the memory is in read mode, and memory read accesses (bytes, 16-bit and 32-bit word) appear like conventional memory. Transitions to read mode can be enabled by writing 0x00ff (little endian) or 0xff00 (big endian) to the base address.

Write mode

Writes are treated as commands, not memory writes. This is where it gets a lot more complicated. The data sheet must be read. Here are some notes.

Two chips – The DE4 has a 512 Mb part containing two 256 Mb dies (chips) in the same package. Therefore, there are actually two independent devices. For example, reading status is on a per-die basis. Address bit 25 determines which die is being used.

Data width – The device is 16 bits wide, and byte-wide accesses make no sense to it. Use only 16-bit writes. 32-bit reads will be turned into two 16-bit reads by our hardware.

Block sizes – Each flash chip is broken down into programming regions and blocks. Blocks are not equal in size. Blocks 0 to 254 are 128 KB in size and blocks 255 to 258 are 32 KB. See Table 7 on page 24 of the data sheet for further details.

Block erase – Data can be erased (set to 0xffff) only by erasing a whole block.

Write protect – After reset, the flash part write-protects the blocks. Software can issue a block unlock request before doing a write, and then lock the block again afterwards. There are also one-time programmable lock registers; we suggest that you avoid touching these!

Writing data – Once a block is unlocked, data can be written one 16-bit word at a time by issuing a write command followed by the data. After doing a write, the status must be polled to determine when the write is complete before another write or read is attempted. Writes can only clear bits; therefore, an erase may be necessary to set all of the bits in the block before doing the write.

Buffered writes – Writes can be conducted in blocks as large as 32×16 -bit words. This is faster than using single writes.

Here are some example access commands (see Table 21, page 51 of the data sheet for further details). Note that this assumes a **little-endian view**:

Read mode (i.e., the same mode as after reset)			
read/write	address	data	comment
write	base address	0x00ff	clear the status register
Unlock block for writes			
read/write	address	data	comment
write	address within block	0x0060	unlock setup
write	address within block	0x00d0	unlock block
Lock block to write protect			
read/write	address	data	comment
write	address within block	0x0060	unlock setup
write	address within block	0x0001	lock block

Status register

Notes on the status register based on Table 28, page 75 of the data sheet:

bit	name	meaning
7	device write status	0=busy, 1=ready
6	erase suspend status	erase suspend 1=not in effect, 0=in effect
5	erase status	0=success, 1=fail
4	program status	0=success, 1=fail
3	V_{pp} status	programming voltage status (0=good, 1=bad)
2	program suspend status	program suspend 1=not in effect, 0=in effect
1	block-locked status	block (0=not) locked during program or erase
0	BEFP status	see data sheet

Bits 7, 6, and 2 are set and cleared by the flash write state machine, but bits 5, 4, 3, and 1 are only set by it. Thus, a clear is needed before using them to check error status.

Note that these tables assume a **little-endian view**:

Clear status register			
read/write	address	data	comment
write	base address	0x0050	clear the status register
Read the status register			
read/write	address	data	comment
write	base address	0x0070	read status register mode
read	base address	-	status register returned

Erase block

Erasing the block requires the following sequence (in pseudo code):

```
unlock_block_for_writes(offset)
clear_status_register
issue_erase_block_command(offset)
while (read_status_register == busy) {} // several million polls
read_status_register to see if erase passed
lock_block_to_prevent_writes
read_mode
```

Note that this table assumes a **little-endian view**:

Erase unlocked block			
read/write	address	data	comment
write	address within block	0x0020	block erase setup
write	address within block	0x00d0	block erase confirm

To erase a region of memory, the easiest approach seems to be to scan the memory to see if it contains 0xffff and, when it does not, issue a block erase command at that address.

Note that Intel certify the part for a minimum of **100,000 erase cycles per block**.

Writes

Write sequence (post erase) starts with an unlock of the block, performs each write followed by a status register check, and finally locks the block again, putting the device back into read mode (as above). The write component is performed using the following sequence (note that this table assumes a **little-endian view**):

Write data			
read/write	address	data	comment
write	address	0x0040	write command
write	address	data	write the data

Then poll the status register (see below) until bit 7 has gone to 1 (ready). This polling typically took 52 polling loop iterations on a NIOS processor running at 100 MHz with each flash access taking 16 clock cycles. (This is not fast!)

offset range	BERI use	device
0x00000000 - 0x00007FFF	user design reset vector	/dev/cfid0s.config
0x00008000 - 0x0000FFFF	ethernet option bits	/dev/cfid0s.config
0x00010000 - 0x00017FFF	board information	/dev/cfid0s.config
0x00018000 - 0x0001FFFF	PFL option bits	/dev/cfid0s.config
0x00020000 - 0x00c1FFFF	FPGA image 1 (power up)	/dev/cfid0s.fpga0
0x00c20000 - 0x0181FFFF	FPGA image 2 (on RE_CONFIGn button)	/dev/cfid1s.fpga1
0x01820000 - 0x03FDFFFF	operating system area	/dev/cfid0s.os
0x02000000 - 0x03FDFFFF	kernel (temporary)	/dev/map/kernel
0x03FE0000 - 0x03FFFFFF	boot loader	/dev/cfid0s.boot

Table 9.5: Layout of the on-board DE4 Intel StrataFlash

Buffered Writes

Buffered writes are more efficient than single writes. The write sequence is only slightly more involved.

Buffered write data			
read/write	address	data	comment
write	address	0x00e8	buffered write command
read	address	status	sr[7]=0 indicates failure
write	address	0x001f	number of data items to write minus one
write	address	data	write 32 words of data
write	address	0x00d0	confirm write
read	address	status	sr[7]=0 means busy, wait

Flash Regions

Terasic specifies uses for most of the flash memory in the *Terasic DE4 Getting Started Guide*. Some of these regions must remain used for their reserved purpose while others have been reallocated for other uses.

In our design, three regions are of particular importance. The region 0x00000000-0x00020000 is reserved for Terasic uses. The region 0x00020000-0x0181FFFF is dedicated to two FPGA images, the first of which is loaded at power-up. This offset is programmed into the MAXII FPGA on the DE4 and cannot be changed easily. The region 0x02000000-0x03FFFFFF (the entire second flash chip) is dedicated to a default software image to be relocated to DRAM at bootup, which is performed by the miniboot bootloader that is embedded in the FPGA image.

Table 9.5 lists our uses for each range and the corresponding FreeBSD device that provides access the region. Changes to these allocations may require changes to this document, miniboot, berictl, BERI.DE4.hints, and flashit.sh.

If portions of the flash are accidentally erased to cause unexpected behavior, factory behavior can be restored by extracting and writing the file `cfi0-de4-terasic.bz` `/dev/cfid0`. This file can be found under `/usr/groups/ctsrtd/cheri` on Cambridge systems.

```
# bunzip -c cfi0-de4-terasic.bz2 | dd of=/dev/cfid0
```

Hardware notes

The device comes out of reset in asynchronous mode operation, which seems to be easiest to deal with. Thus, the clock to the flash device is simply kept at 0.

The bus is simple to use. Address, address-valid, chip-enable, write-enable, and output-enable can be asserted together. Writes take a minimum of 85 ns and the address and data are latched on the rising edge of write-enable. The choice to deassert chip-enable (i.e., set to 1) between each access seems to guarantee updates.