


CheriOS

Talk by: Lawrence Esswood

CHERI (briefly)

- Capabilities
 - Identify resources & grants right to use them
 - Unforgeable
- CHERI
 - Hardware enforced capabilities for
 - Memory (split into read/write/execute...)
 - System management
 - Sealing...

A CHERI Capability



A CHERI Capability

- Not just for bounds checking
 - Compartments
 - Sealing

Principles

- CheriOS, designed with failure in mind.
 - Unreasonable to expect people to build systems that are both correct and large

Principles

- Compartmentalise
 - Limit exposure on failure

Principles

- Least Privilege
 - No point splitting roles if each actor has the same privilege

Principles

- Privilege does not imply trust
 - Assume privilege will always be misused

Principles

- Failure should result in only denial of service
 - A system component can
 - Fail to provide its service
 - Provide the service incorrectly
 - A system component cannot
 - Compromise the users confidentiality/integrity
 - Confuse the user into thinking they have been provided a resource they have not

Examples - Scheduler

- Purpose
 - Provide time sharing
- Can
 - Deny scheduling you completely
 - Wake you up from sleeping too early
- Cannot
 - Access your state
 - Change your state when it reschedules you
 - Convince you an event has happened when it hasn't

Examples – Page Allocator

- Purpose
 - Provide memory resources to users
- Can
 - Fail to provide you memory
 - Take back memory unexpectedly (causing a fault)
- Cannot
 - Read the memory it gives you
 - Change the contents unexpectedly

One Ring to Rule Them All

- Privilege Rings have problems
 - Inherently hierarchical
 - Create high-value targets (get at a given ring, and every outer ring is yours)
 - Large TCBs
 - A user is forced to trust the entire OS
 - e.g. Windows has 45 MLOC!
- Can we build a **flat** system?
- Can we build an **untrusted** system?

Fundamental Problems

- System interfaces are too powerful
 - Exception Vectors
 - Page Tables
- CHERI capability derivation is inherently hierarchical
 - How can a user gain a capability, via the OS, that the OS does not have?
 - How can we stop the user from abusing such a mechanism if it did exist?

The Nanokernel

- Division of system capabilities into larger, less powerful, orthogonal set
- Provide new security primitives
 - Reservations (Memory Integrity & Confidentiality)
 - Foundations (Program Identity & Attestation)
- Really Small!
 - 2625 MIPS **instructions** (about 50 routines)
 - No stack, nearly entirely loop free preemptable leaf functions
 - Compare to microkernel like SEL4, about 10 000 Lines of C

CheriOS

- Microkernel (pre-emptable)
 - Scheduler activations, Message Passing, Interrupts/Exceptions
- All other OS services fully compartmentalised
 - Memory Manager
 - Process Manager
 - WARNING
 - Namespace Service
 - File Systems / Drivers / Type manager ...
- All 'unprivileged'
 - Apart from having nanokernel capabilities a user may not

What can a user expect?

- Spatial Safety
- Temporal Safety
- Compartmentalisation
 - Only explicit argument capabilities can move outside compartment
- Compartment local CFI

The Nanokernel – System Management

- Opaque VCPU contexts
 - Capability to create / destroy
 - Capability to switch
 - Capability to set a context as the exception context
- (Non user) exceptions force context switch
 - Exception context cannot modify state of victim context
- User exceptions can be handled by compartments
 - 210 / 387 (to assembly / to C) cycles to handler and back
- NOT a compartment
 - Nanokernel mostly agnostic to compartments

The Nanokernel – System Management

- Checks Page Table Operations
 - Presently enforces bijective mapping
 - Cannot re-map without either
 - Revoking
 - Proving this would grant no extra power (future extension)
 - Only allows mapping to pages that have been zeroed
 - This makes all memory (by default)
 - Start in a guaranteed state
 - Be non aliasing
 - Seem like one big physical memory, or gives a fault

Reservations

- Capability to *allocate* a range of memory
- Can be copied
- All operations are (somewhat) destructive and apply to all copies
- Software enforced

Reservations - State

- Open
 - Available for allocation (no client has access)
- Taken
 - Has been allocated (allocating client has access)
 - Handle used for revoking
- Merged
 - Revoking delegated to another handle

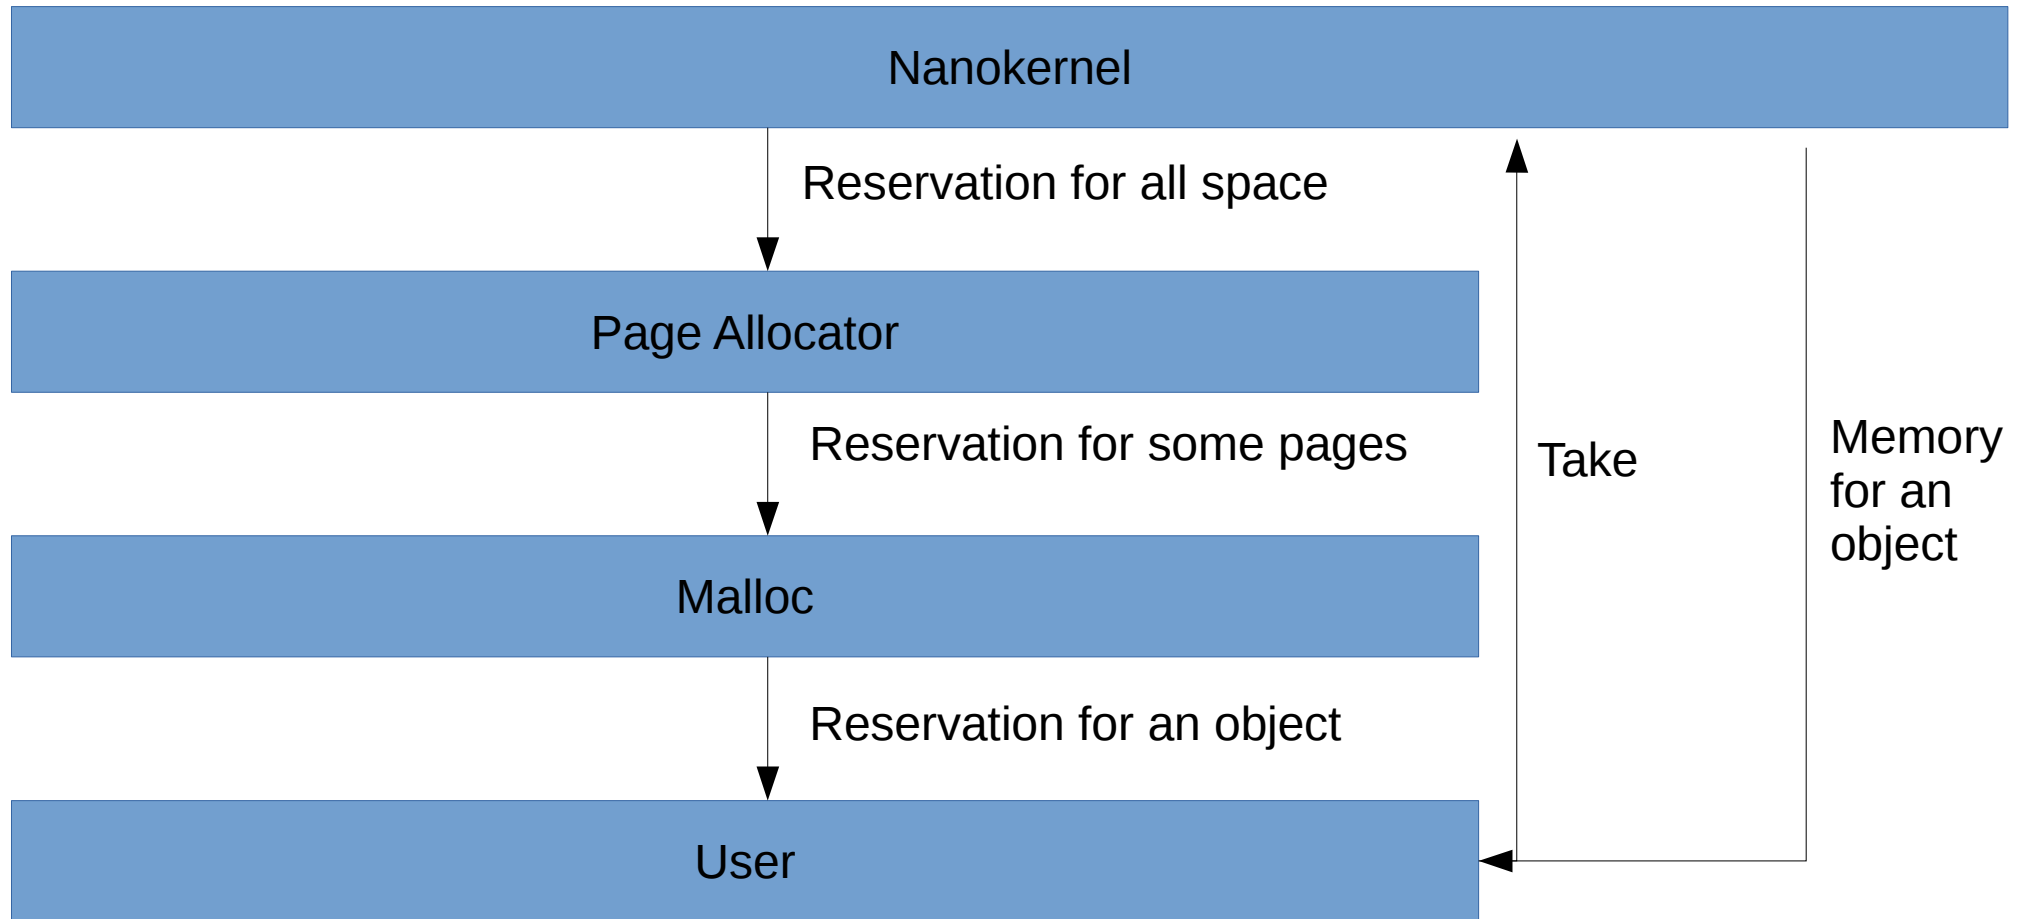
Reservations - Operations

- Take
 - Open → Taken
- Split
 - Open → (Open, Open)
- Merge
 - (Taken, Taken) → (Taken, Merged)
- Revoke
 - Taken → Revoking → Open
 - Destroys all capabilities to region
 - Destroys ALL stale handles (including merged ones)

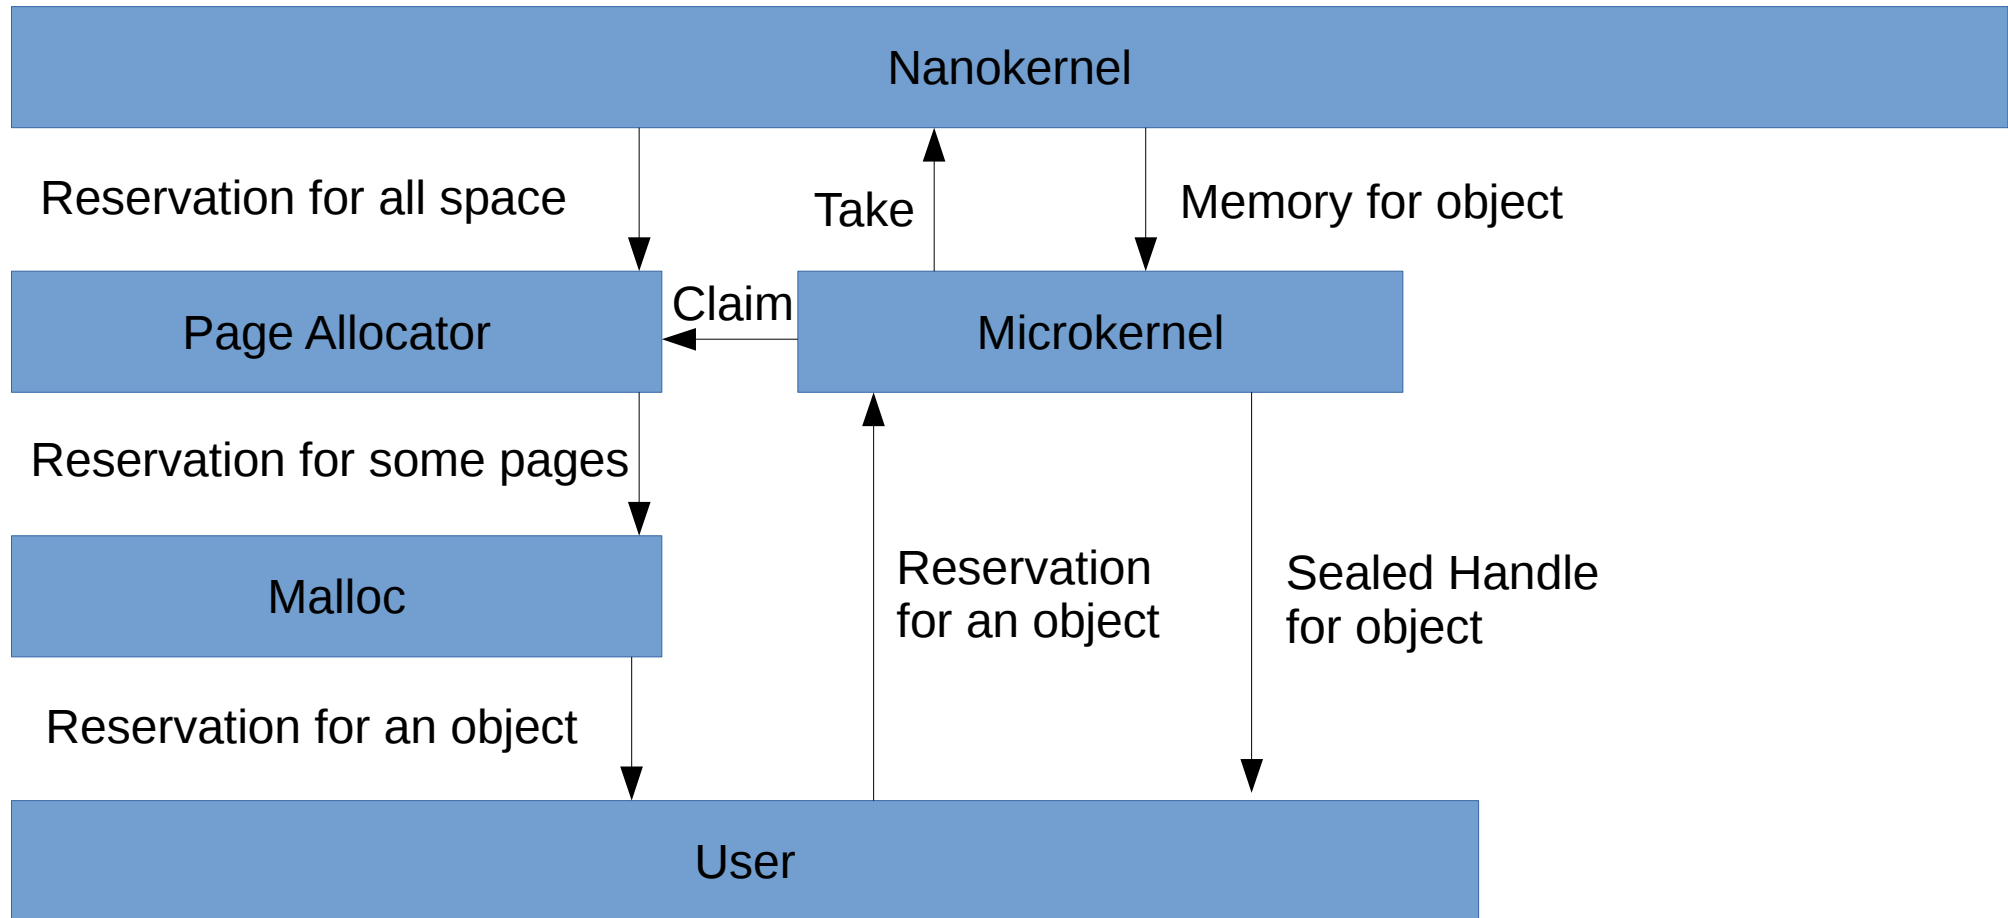
Reservation Guarantees

- When you take a reservation
 - You get a memory capability
 - It is unique and non-aliasing (outside the nanokernel)
 - The nano kernel will not, even in part, grant the same capability again
 - Memory will be in a known starting state (all zero)
- This relies on the rules for virtual memory

Reservations – Typical Allocation



Reservations – More interesting



Spatial Safety

- Get reservations from malloc
 - Guaranteed non aliasing allocation
 - Don't trust malloc
 - Don't trust the OS
 - Low overhead of few extra function calls to allocate
 - **No** overhead when being used
- Exploit CHERI bounds checkis
 - All objects have strict bounds checking
 - Including function bodies

Temporal Safety

- Non re-use
 - Allocations cannot alias with old ones
 - Can re-use/pool if performance is critical
- For the stack as well (selectively)!
 - New 'unsafe' keyword for C
 - Unsafe will not re-use that stack space
 - Automatic compiler assistance for safe/unsafe tagging
 - Less costly than using malloc for every stack frame
 - No cost if keyword is not used and allocations can be proven safe
 - True for majority of frames both statically and dynamically

Compartmentalisation

- Every dynamic library (or even function) can be its own sandbox
 - Elective, potentially uni-directional isolation
- New calling convention
 - Callers can make an 'untrusting' call
 - Callees can elect for Callers to be 'Untrusted'
 - Cost only paid if used
 - Dynamically reconfigurable
 - The microkernel/nanokernel are just dynamic libraries!
 - Syscalls/nanocalls have the same cost as dynamic library calls

CHERI (sealing/ccall)

- Capability to 'seal/unseal'
 - With a particular type
- Sealed capabilities cannot be used
- CCall
 - Need two capabilities, one code, one data, of the same type
 - Unseals both and atomically jumps to the code
 - Puts data in IDC
 - In Cheros IDC always points to a structure that you should think of as a compartment identifier

Compartment

- Not a VCPU / Scheduler activation
- Identified by a data structure (DLS) used for ccall
 - Per thread
 - Can be used from any VCPU / Scheduler activation
- DLS has entries for (depending on subtype)
 - Subtype
 - Concurrency guard
 - User exception handling
 - Sealing / Unsealing capability
 - Stack
 - Capability to program globals table
 - Capabilities to thread local variables

Calling Convention

- Untrusting call (callee untrusted)
 - Save all registers (even callee saves)
 - Make and seal a new DLS
 - Zero registers
 - On return make CFI checks and restore
- Untrusted call (caller untrusted)
 - Make CFI checks on entry
 - On return clear tmps and caller saves

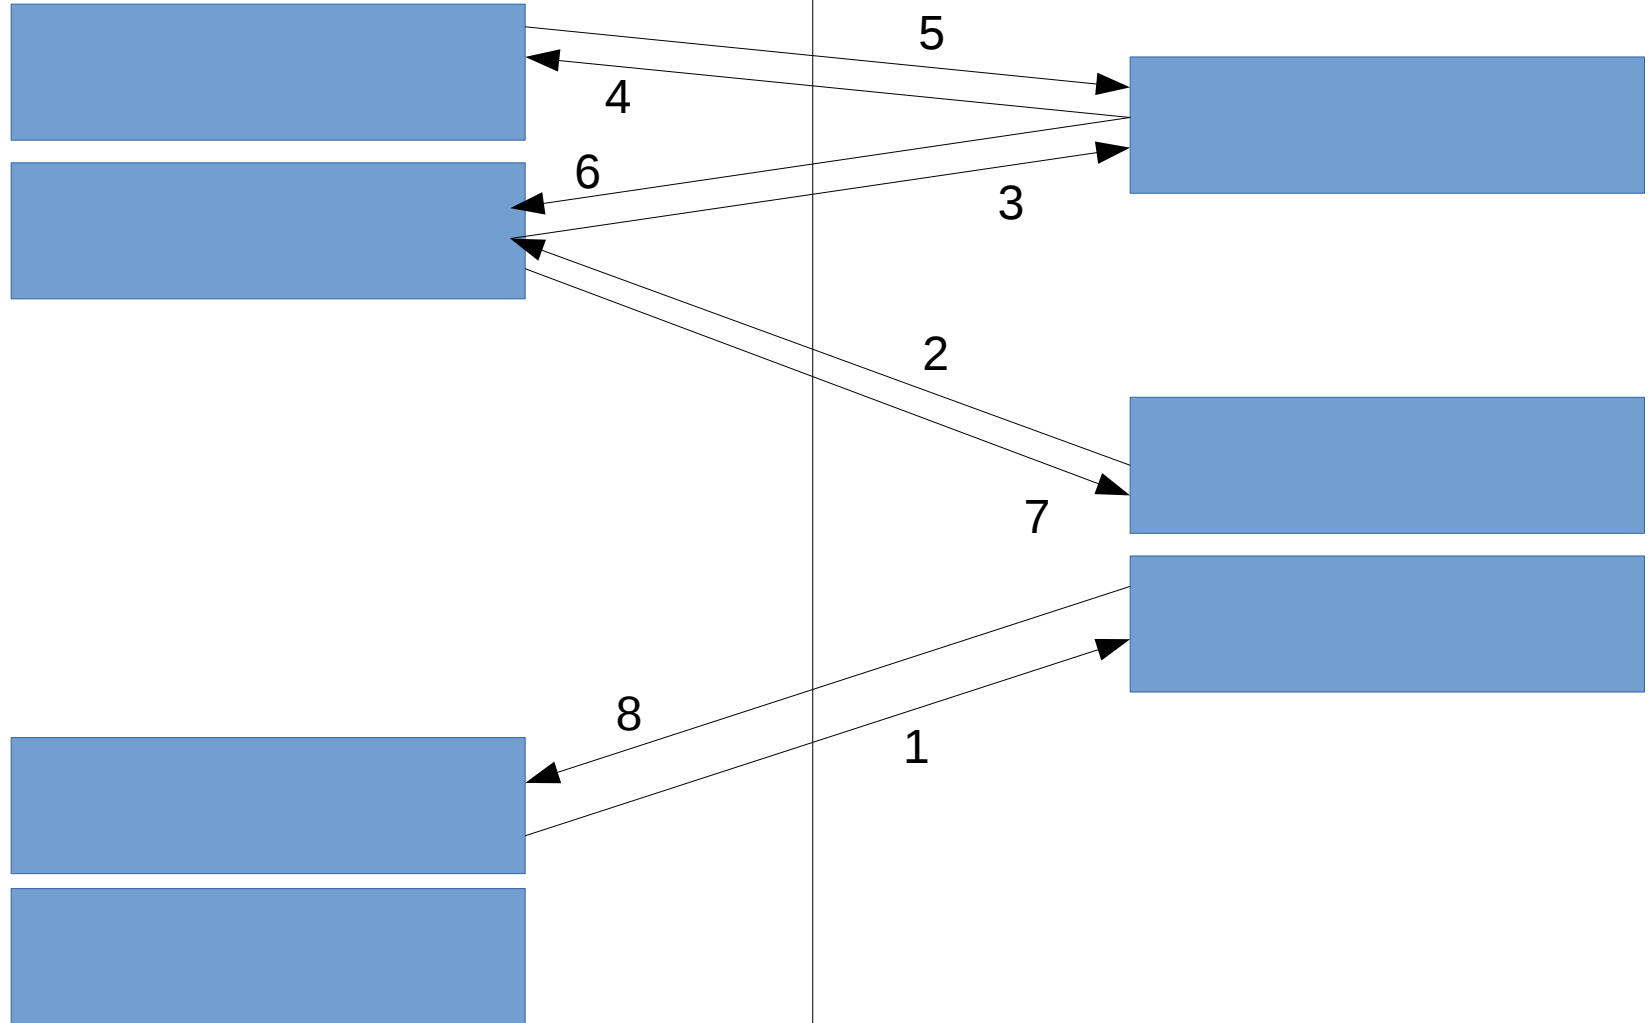
CFI

- CFI edges enforced between compartments
- Calls:
 - Cannot call into a compartments with a particular C thread if already in compartment
 - Can only call starts of functions, and only if exported
 - Cannot return to compartment out of order
 - Cannot return to compartment if return has already been used

CFI

My Compartment


Other Compartment



CFI

My Compartment

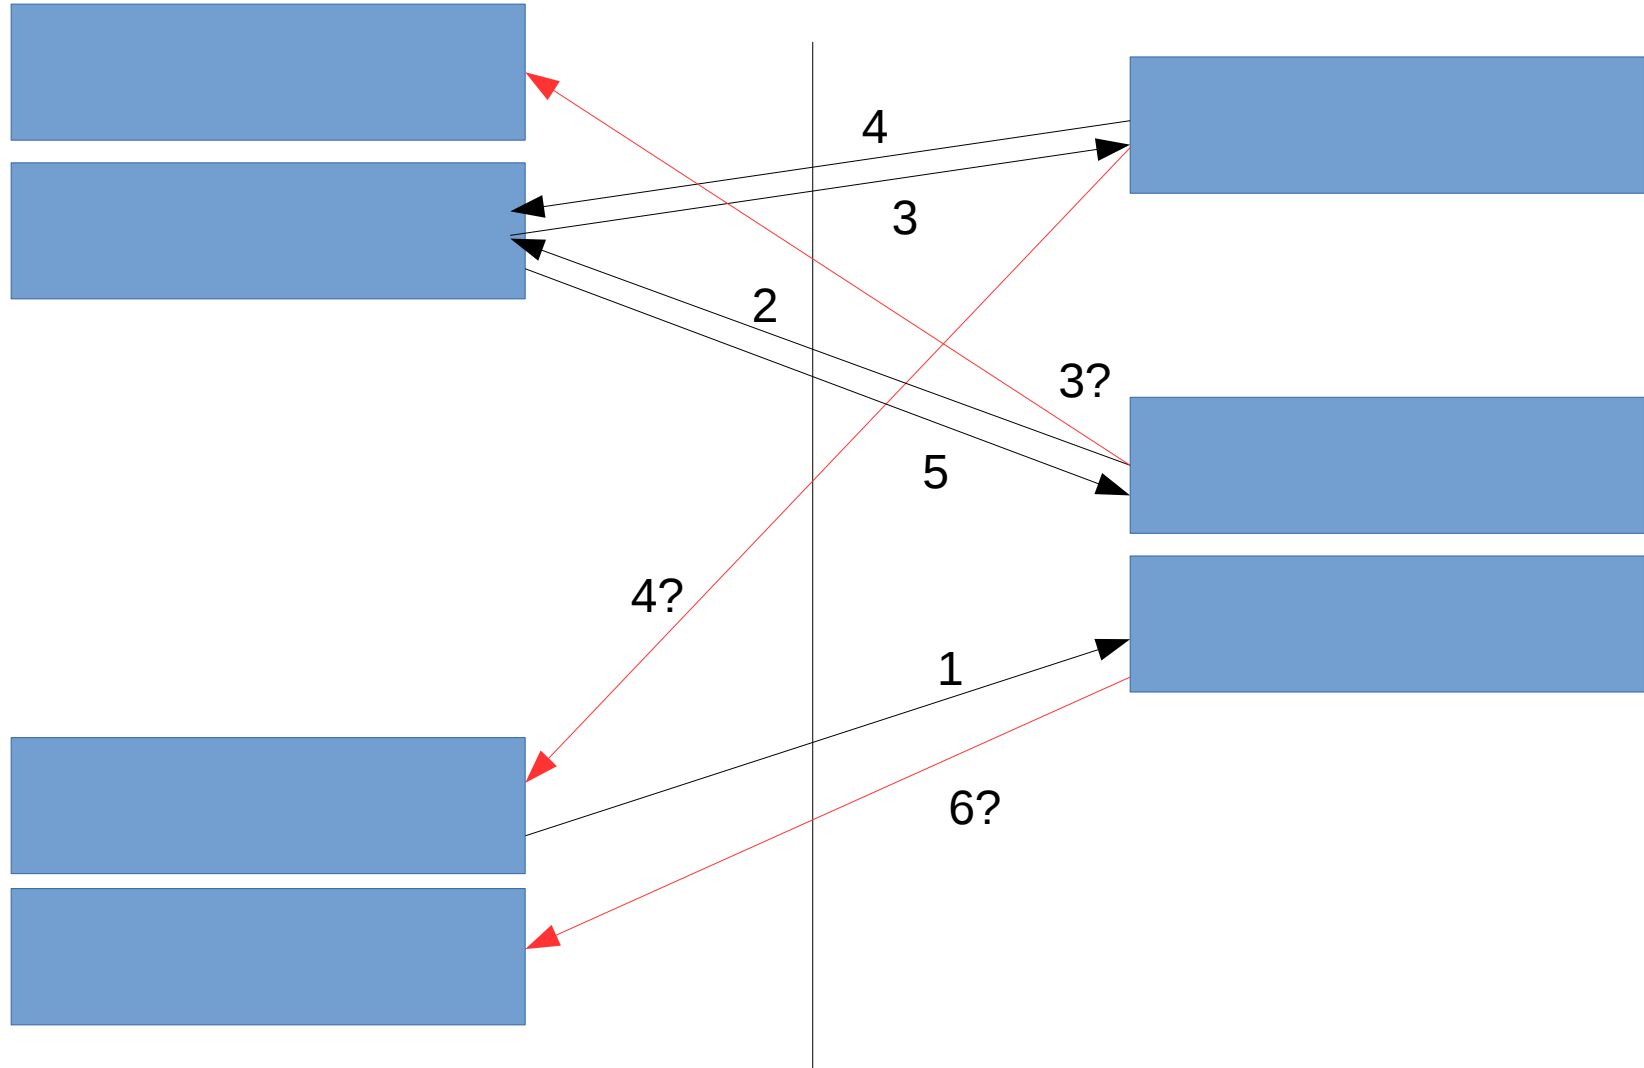
Other Compartment



CFI


My Compartment

Other Compartment




CFI

My Compartment



Other Compartment




Sharing Capabilities

- Users are encouraged to share capabilities directly
 - Even across process boundaries!
- Means we cannot ever reuse memory
 - Temporal safety ensures we do not
- Lifetime gets more complicated
 - Malloc augmented with 'claim'
 - Magic safe dereference instruction
 - Or just use exceptions


Foundations

- Identity
 - How we take a digest of a capability graph?
- Attestation
 - Given we can trust our own execution, can we verify another program has been loaded correctly?
- Send verifiable messages from and to

Capability Graph - Digest




Foundations - Digest



Foundations - Creation

```
entry_t foundation_create(res_t res, size_t image_size,  
                           capability image, size_t entry0,  
                           size_t n_entries, register_t is_public)
```



Foundations

- Each foundation created has an identity
- Identity is a read-only memory capability
 - A sha256 of some contiguous bytes
 - A length
 - An offset of the first entry point

Foundations - Entry

```
void foundation_enter(entry_t entry)
```

- Jumps to entry point specified by entry token
- Creates a R/W capability for the foundation
- Grants an 'authority' token

Foundations - Exit

- Exit is implicit
 - Lose access to the capability
 - Already done by untrusting call
- Re-entry is also implicit
 - Explicit call only generally used for setup
 - Entry / Exit therefore has no overhead over normal untrusting call

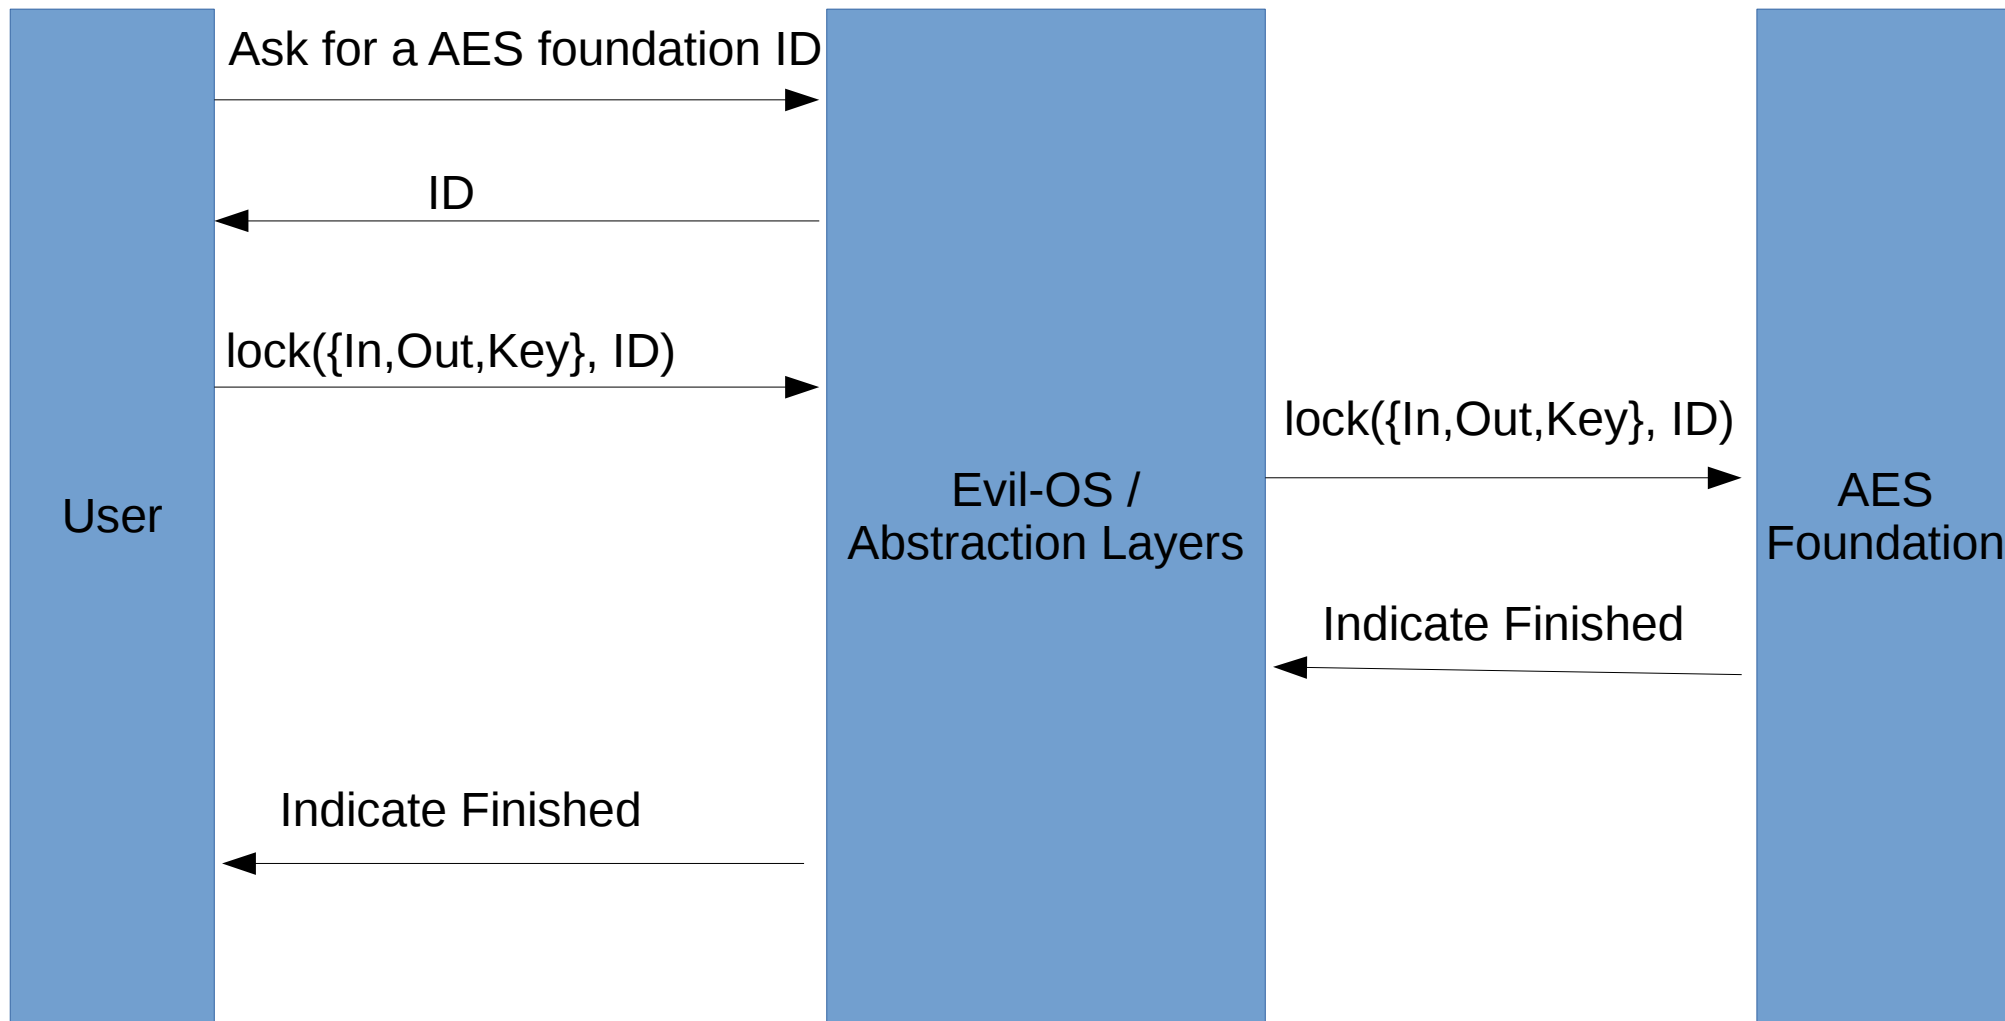
Foundations

- Authority token
 - Symmetric Key, can lock-sym unlock-sym
 - Asymmetric Key, can sign and unlock-asym
 - Can also derive identity token
- Identity Token
 - Asymmetric Public Key, can lock-asym

Foundations

- We can now treat the operating system and IO as an untrusted channel
- Foundations are identities, not compartments
 - i.e., multiple compartments may have the same or multiple authority tokens
- Replace encryption with foundation primitives

Example: AES



Foundation Uses

- Load an entire OS in a foundation
 - Secure VM boot
- Load an entire program
 - Secure against malicious program loader / OS
- Secure Deduplication
- Callback signing

Foundations vs encryption

- Scalable
 - Merkle trees do not scale
 - Encryption is expensive
- Nestable
 - Can hold multiple, not necessarily hierarchical, set of authority tokens

Revocation

- Capabilities get *everywhere*

Revocation - Stage1

- Unmap
 - Nano kernel will not allow remapping virtual
 - Nano kernel will not allow remapping physical until zeroed
 - Allows for ~60bits of space before we cannot use this method
- Can merge old reservation handles
 - Only need one handle per finished range
 - Can unmap metadata as well!

Revocation – Stage2

- Concurrent scan of all **physical** memory
 - Revoke a vast **virtual** range
 - Can do without suspending any users
 - Takes on the order of ms
- Will return a new open reservation for the virtual space

Problems with CheriOS

- This is all defunct with DMA
 - Need a capability aware DMA engine
- Fragmentation due to non-reuse (heap)
 - Need nano kernel assisted compacting allocator

State of CheriOS

- Runs on QEMU/FPGA CHERI-MIPS
- Has a FAT file-system
- Runs a HTTP server
 - Webstack is a port of LWIP
 - HTTP server is a port of NGINX
 - Hosting a webpage at cherios-fpga.sec.cl.cam.ac.uk
- Publicly available!
 - <https://github.com/CTSRD-CHERI/cherios/tree/lawrence>

FPGA Numbers

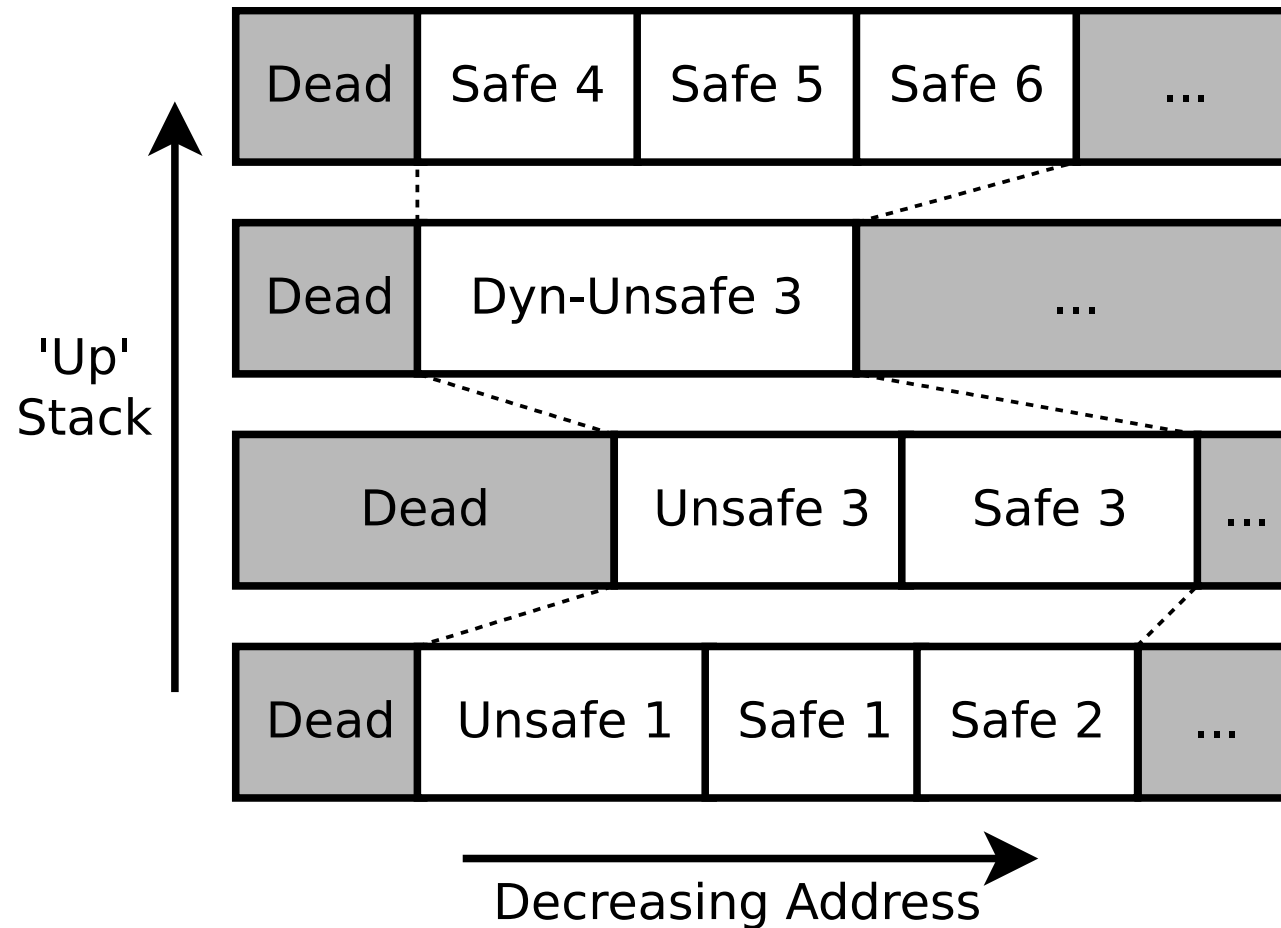
- RPC (Cross process microkernel sync message-send, 2 context switches)
 - 1056 cycles (Trust microkernel)
 - 1256 cycles (Distrust microkernel)
- Calling between compartments (call and return)
 - 5 instructions, 26 cycles (To nanokernel)
 - 71 cycles, Trusting / Trusted
 - 119 cycles, Trusting / Untrusted
 - 152 cycles, Untrusting / Trusted
 - 210 cycles, Untrusting / Untrusted

Questions!

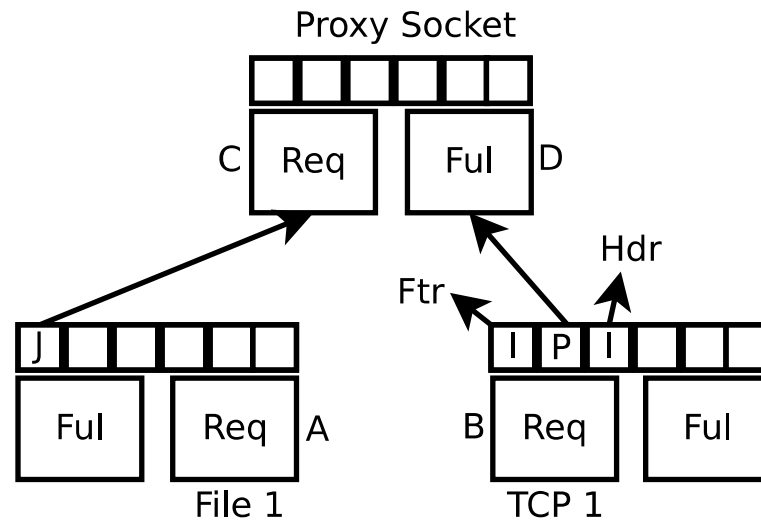

Public foundation

- Contents are not secret, but should not be modified
- Can request read only capability given an ID
- Used for secure deduplication

Slinky Stacks



CheriOS sockets



The Nanokernel - Access

- Accessed via CCall
- Interface can be provided by syscall
 - syscall will always trap to the nanokernel
 - Nano kernel will not deliver this syscall to the system
 - Allows users to circumvent the OS if they desire
 - OS can tell nano kernel to not give out particular interfaces to users